

Towards a unified treatment of the dynamics of collective learning

Wolfram Barfuss

Challenges and Opportunities for Multi-Agent Reinforcement Learning - COMARL AAAI 2021

March 22-23, 2021



Collective Learning Dynamics

An unusual presentation - no results, only questions

Collective Learning Dynamics

An unusual presentation - no results, only questions

- What do I mean with Collective Learning Dynamics?

Collective Learning Dynamics

An unusual presentation - no results, only questions

- What do I mean with Collective Learning Dynamics?
- Multiple literatures → **can we learn from each other?**

Collective Learning Dynamics

An unusual presentation - no results, only questions

- What do I mean with Collective Learning Dynamics?
- Multiple literatures → **can we learn from each other?**
- **How to use our (unified) knowledge**
to help solving challenges of multi-agent reinforcement learning?

Collective Learning Dynamics

An unusual presentation - no results, only questions

- What do I mean with Collective Learning Dynamics?
- Multiple literatures → **can we learn from each other?**
- **How to use our (unified) knowledge**
to help solving challenges of multi-agent reinforcement learning?



My wish

Follow-up interactions
& collaborations

Learning dynamics?

Learning dynamics?

=> **Replicator Reinforcement Learning Dynamics**

based on the link between evolutionary game theory and reinforcement learning

Learning dynamics?

=> **Replicator Reinforcement Learning Dynamics**

based on the link between evolutionary game theory and reinforcement learning

Stochastic RL algorithms



= Idealised learning equations

Learning dynamics?

=> **Replicator Reinforcement Learning Dynamics**

based on the link between evolutionary game theory and reinforcement learning

Stochastic RL algorithms



= Idealised learning equations

Example
Emergence of cooperation

Learning dynamics?

=> **Replicator Reinforcement Learning Dynamics**

based on the link between evolutionary game theory and reinforcement learning

Stochastic RL algorithms



= Idealised learning equations

Example

Emergence of cooperation

Prisoner's Dilemma

A tragedy of the commons

		Agent 1
		Co-operate Defect
Agent 2	Co-operate	1, 1 -2, 3
	Defect	3, -2 0, 0

Learning dynamics?

=> **Replicator Reinforcement Learning Dynamics**

based on the link between evolutionary game theory and reinforcement learning

Stochastic RL algorithms



= Idealised learning equations

Example

Emergence of cooperation

Prisoner's Dilemma

A tragedy of the commons

		Agent 1	
		Co-operate	Defect
Agent 2	C	1, 1	-2, 3
	D	3, -2	0, 0

Learning dynamics?

=> **Replicator Reinforcement Learning Dynamics**

based on the link between evolutionary game theory and reinforcement learning

Stochastic RL algorithms



= Idealised learning equations

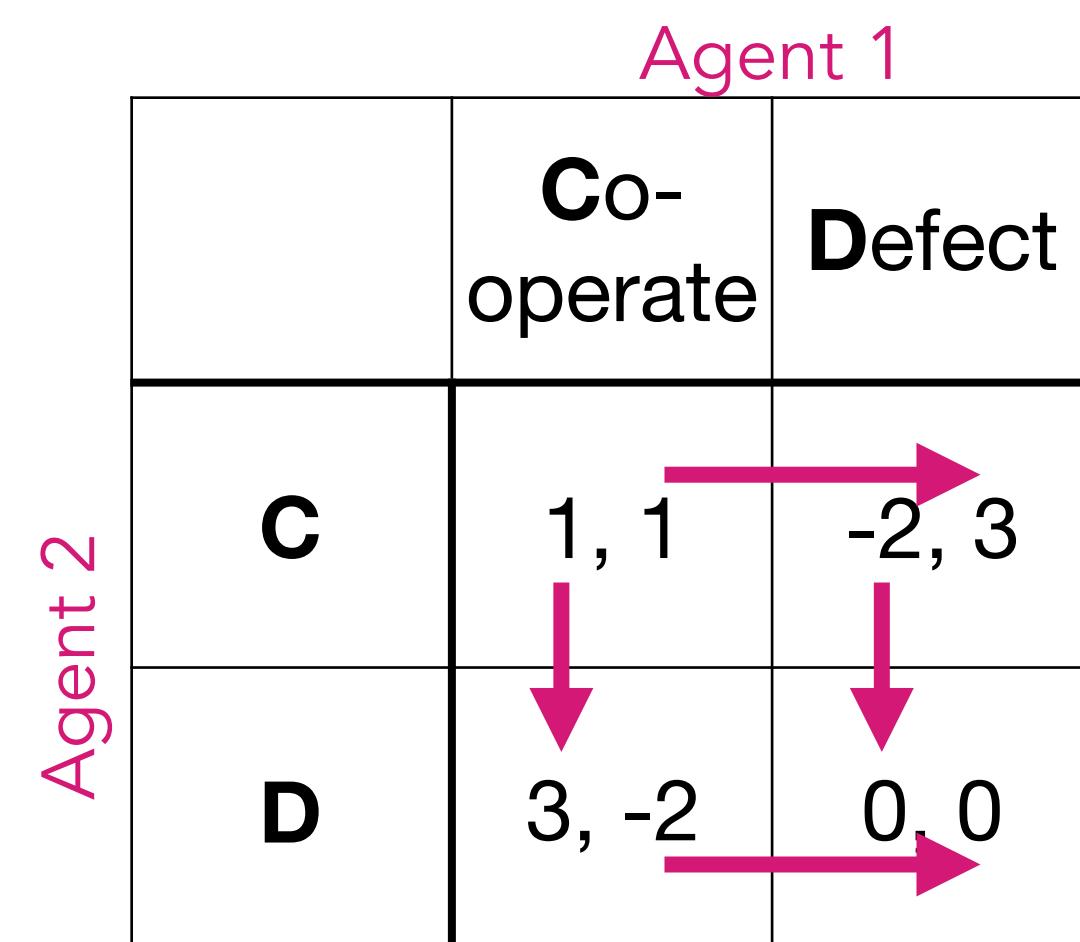
Example

Emergence of cooperation

Prisoner's Dilemma

A tragedy of the commons

		Agent 1	
		Co-operate	Defect
Agent 2	C	1, 1	-2, 3
	D	3, -2	0, 0



The diagram shows a 2x2 matrix representing the Prisoner's Dilemma. The columns are labeled 'Agent 1' and the rows are labeled 'Agent 2'. The strategies are 'Co-operate' and 'Defect' for Agent 1, and 'C' and 'D' for Agent 2. The payoffs are: (C, C) = (1, 1), (C, D) = (-2, 3), (D, C) = (3, -2), and (D, D) = (0, 0). Arrows point from the (C, C) and (D, D) cells to the (C, D) and (D, C) cells respectively, indicating transitions between strategies.

Learning dynamics?

=> **Replicator Reinforcement Learning Dynamics**

based on the link between evolutionary game theory and reinforcement learning

Stochastic RL algorithms



= Idealised learning equations

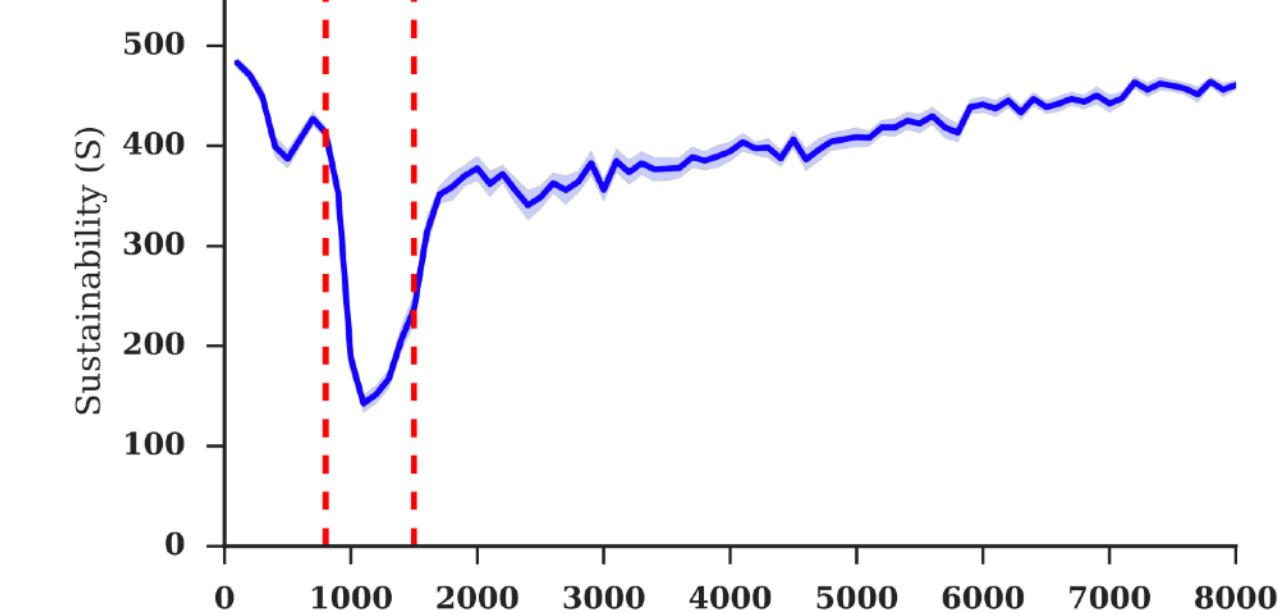
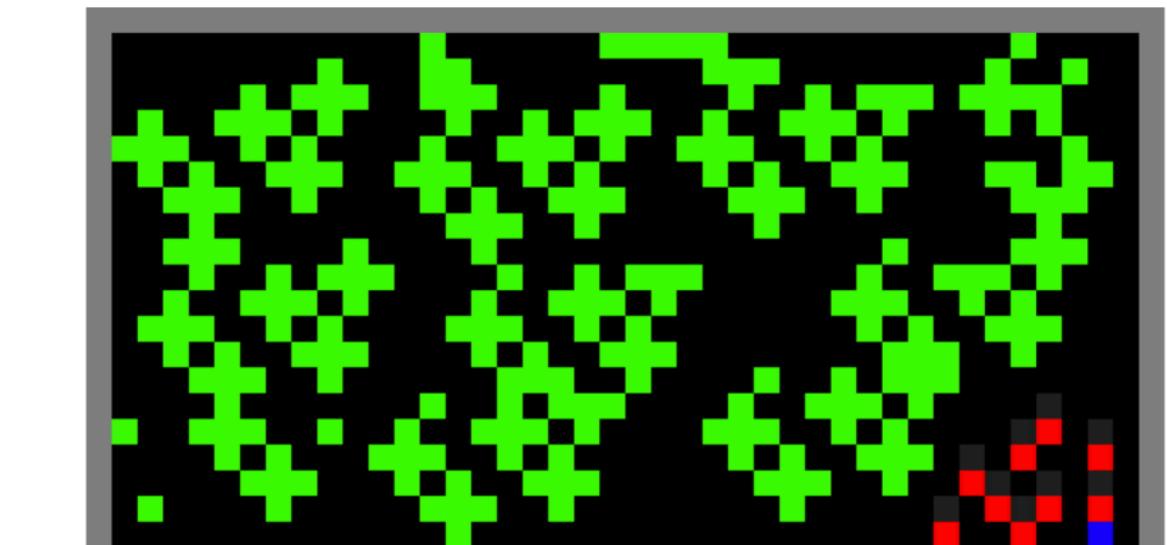
Example Emergence of cooperation

Prisoner's Dilemma

A tragedy of the commons

		Agent 1	
		Co-operate	Defect
Agent 2	C	1, 1	-2, 3
	D	3, -2	0, 0

Deep MARL for common-pool resource appropriation
Perolat et al, 2017



Learning dynamics?

=> **Replicator Reinforcement Learning Dynamics**

based on the link between evolutionary game theory and reinforcement learning

Stochastic RL algorithms



= Idealised learning equations

Example Emergence of cooperation

Prisoner's Dilemma

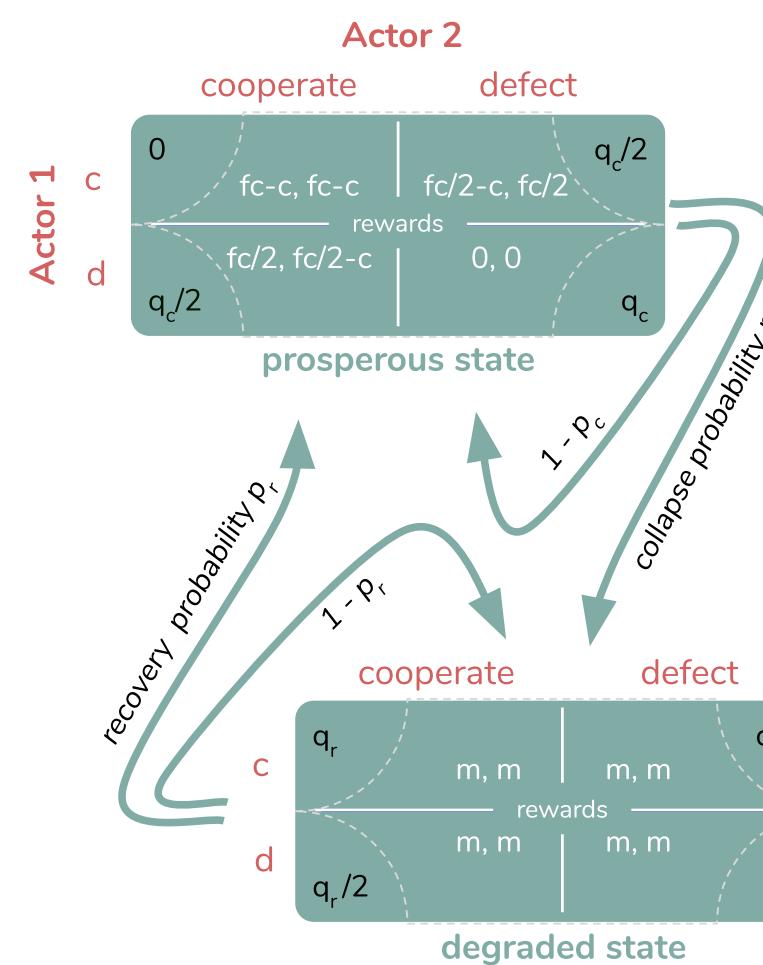
A tragedy of the commons

		Agent 1
		Co-operate Defect
Agent 2	Co-operate	1, 1 -2, 3
	Defect	3, -2 0, 0



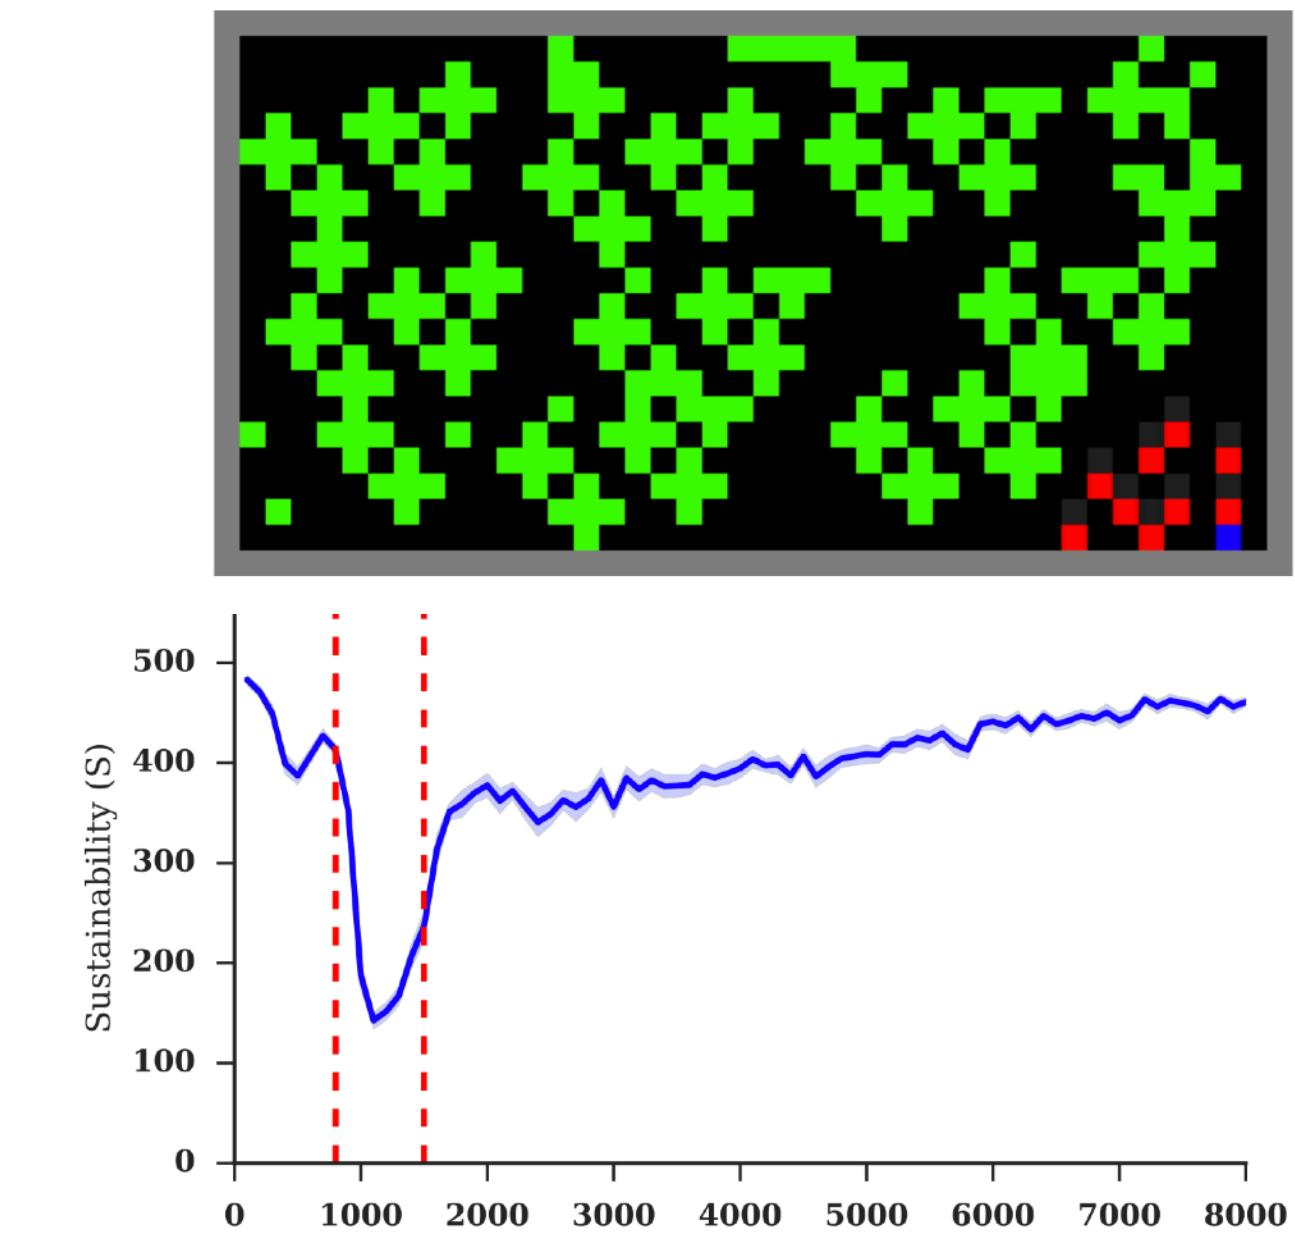
Social-ecological dilemma

Barfuss et al., 2020



Deep MARL for common-pool resource appropriation

Perolat et al, 2017



Learning dynamics?

=> **Replicator Reinforcement Learning Dynamics**

based on the link between evolutionary game theory and reinforcement learning

Stochastic RL algorithms



= Idealised learning equations

Example Emergence of cooperation

Prisoner's Dilemma

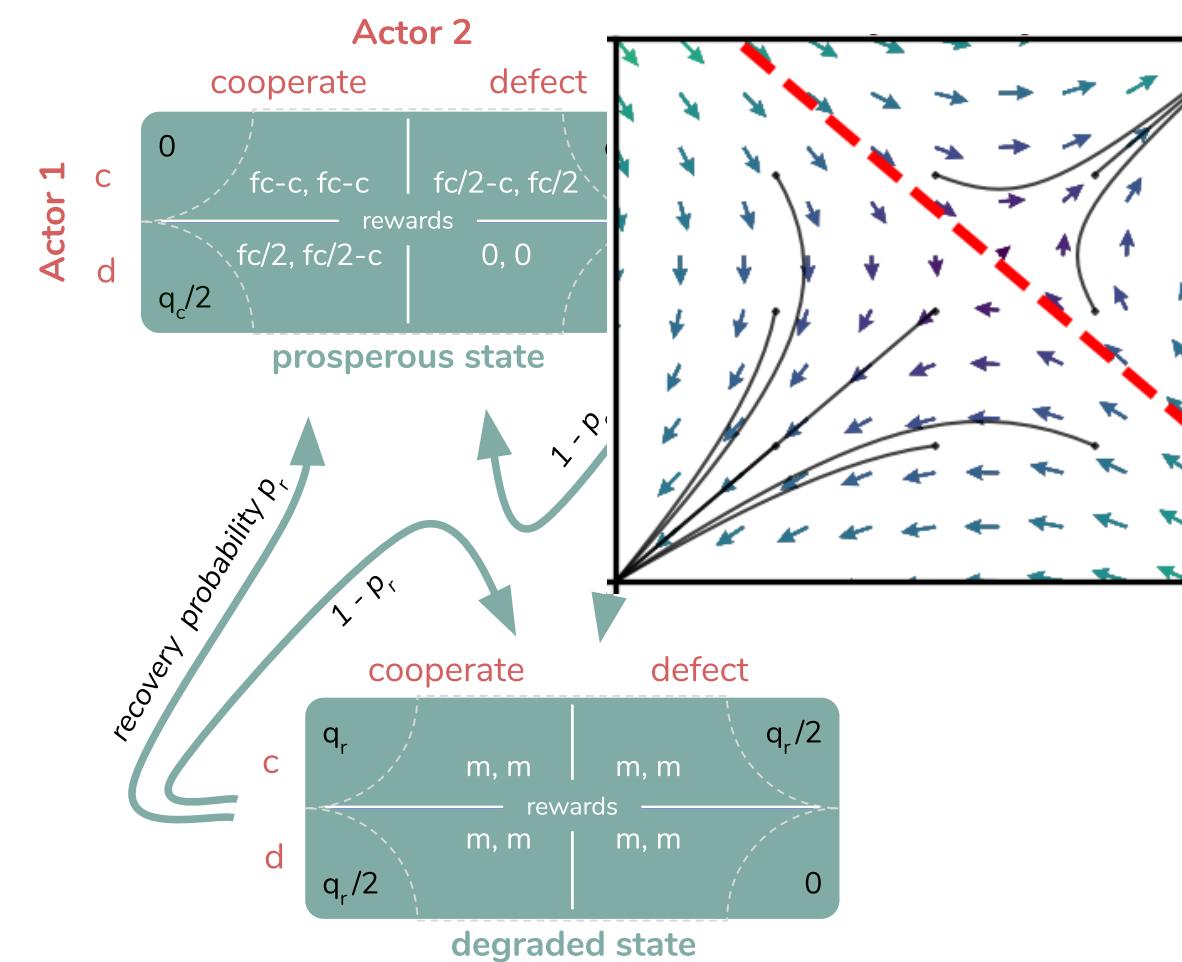
A tragedy of the commons

		Co-operate	Defect
		Co-operate	Defect
Agent 1	C	1, 1	-2, 3
	D	3, -2	0, 0



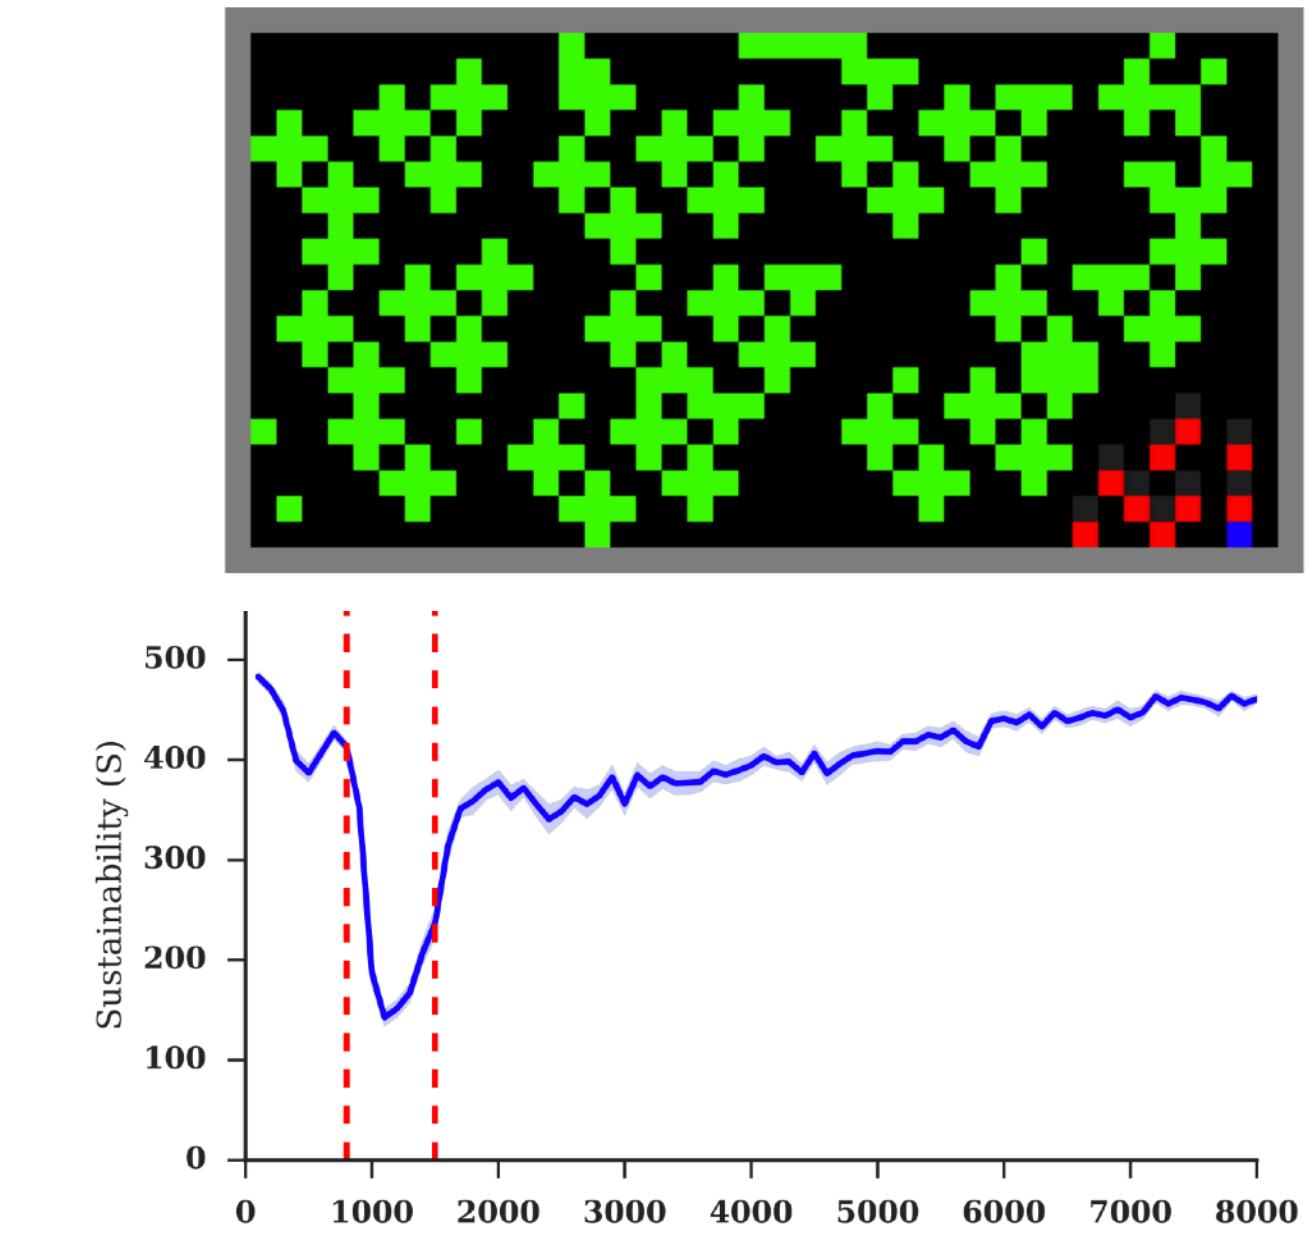
Social-ecological dilemma

Barfuss et al., 2020



Deep MARL for common-pool resource appropriation

Perolat et al, 2017



Learning dynamics?

=> **Replicator Reinforcement Learning Dynamics**

based on the link between evolutionary game theory and reinforcement learning

Stochastic RL algorithms



= Idealised learning equations

Example Emergence of cooperation

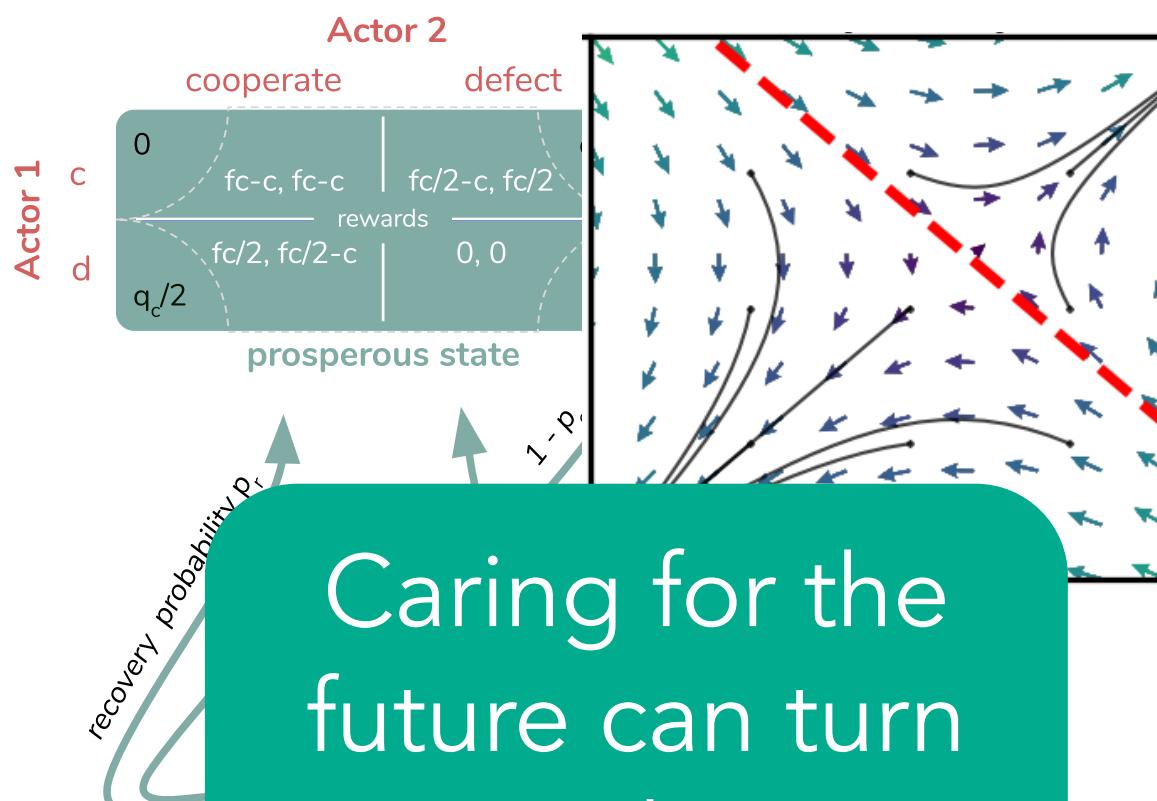
Prisoner's Dilemma

A tragedy of the commons

		Agent 1	
		Co-operate	Defect
Agent 2	C	1, 1	-2, 3
	D	3, -2	0, 0

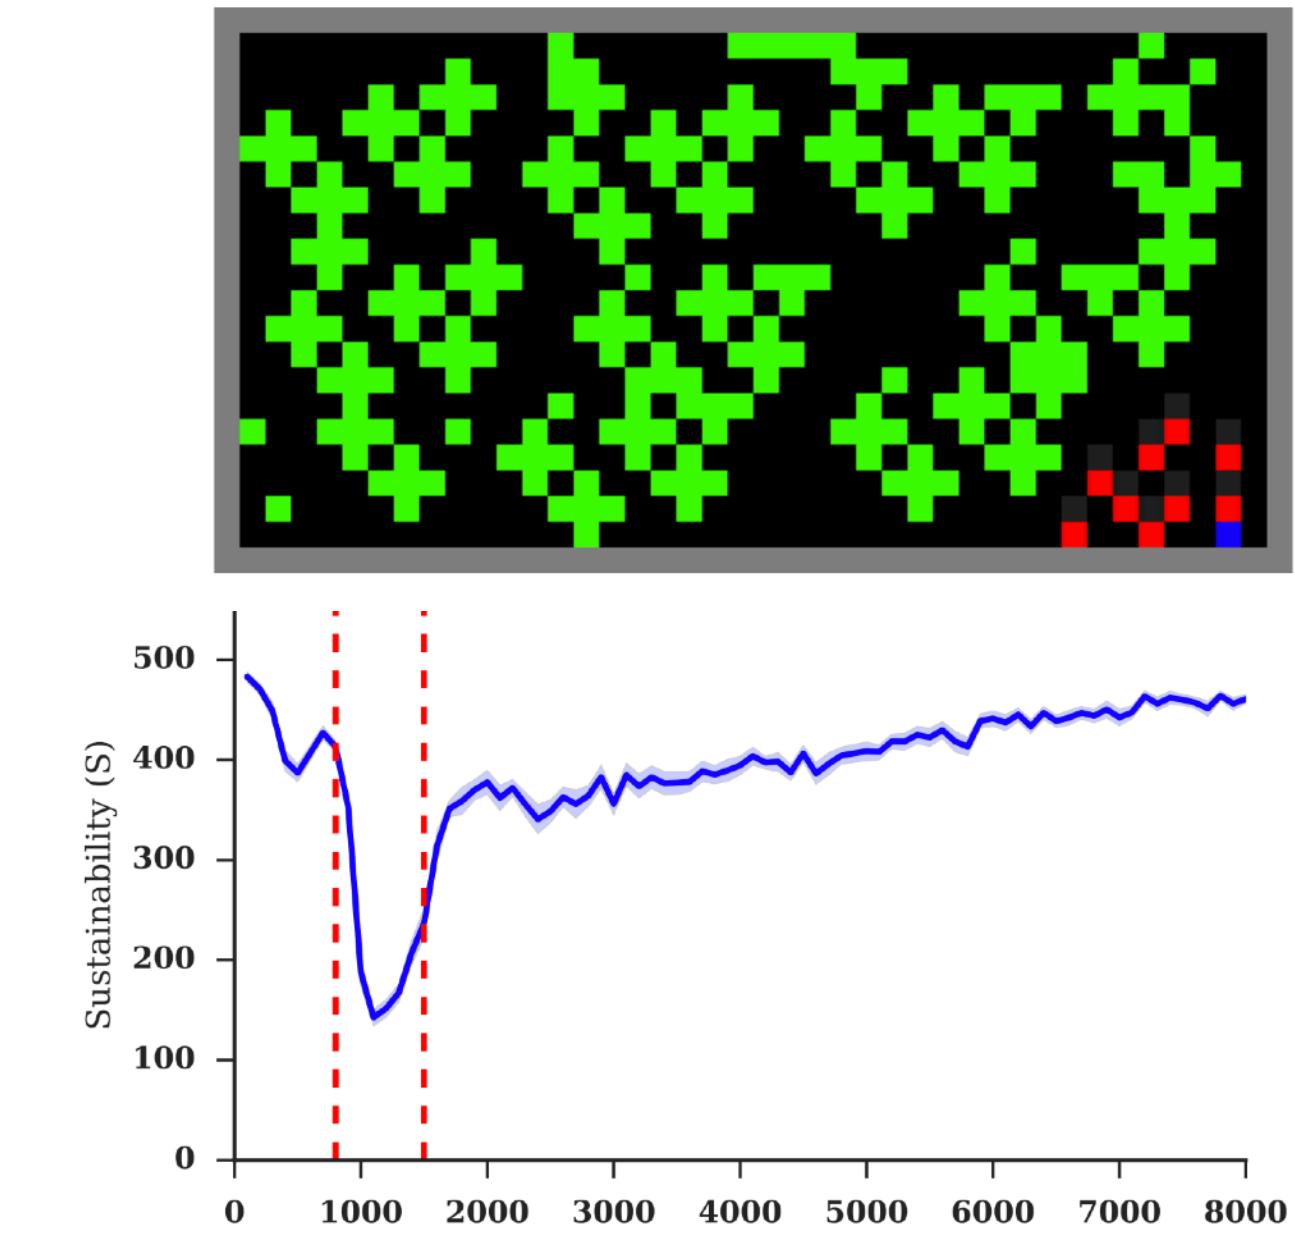
Social-ecological dilemma

Barfuss et al., 2020



Deep MARL for common-pool resource appropriation

Perolat et al, 2017



Learning dynamics?

=> **Replicator Reinforcement Learning Dynamics**

based on the link between evolutionary game theory and reinforcement learning

Stochastic RL algorithms



= Idealised learning equations

Example Emergence of cooperation

Prisoner's Dilemma

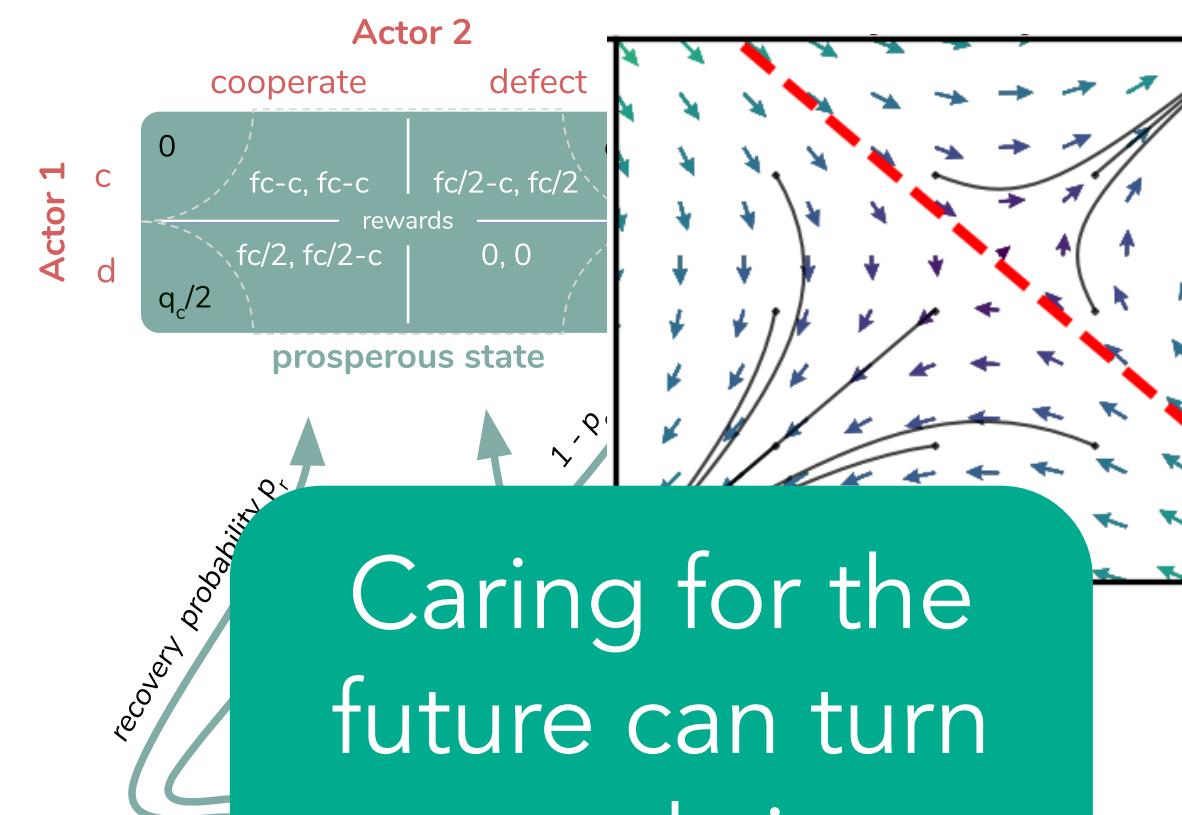
A tragedy of the commons

		Agent 1
		Co-operate Defect
Agent 2	Co-operate	1, 1 -2, 3
	Defect	3, -2 0, 0



Social-ecological dilemma

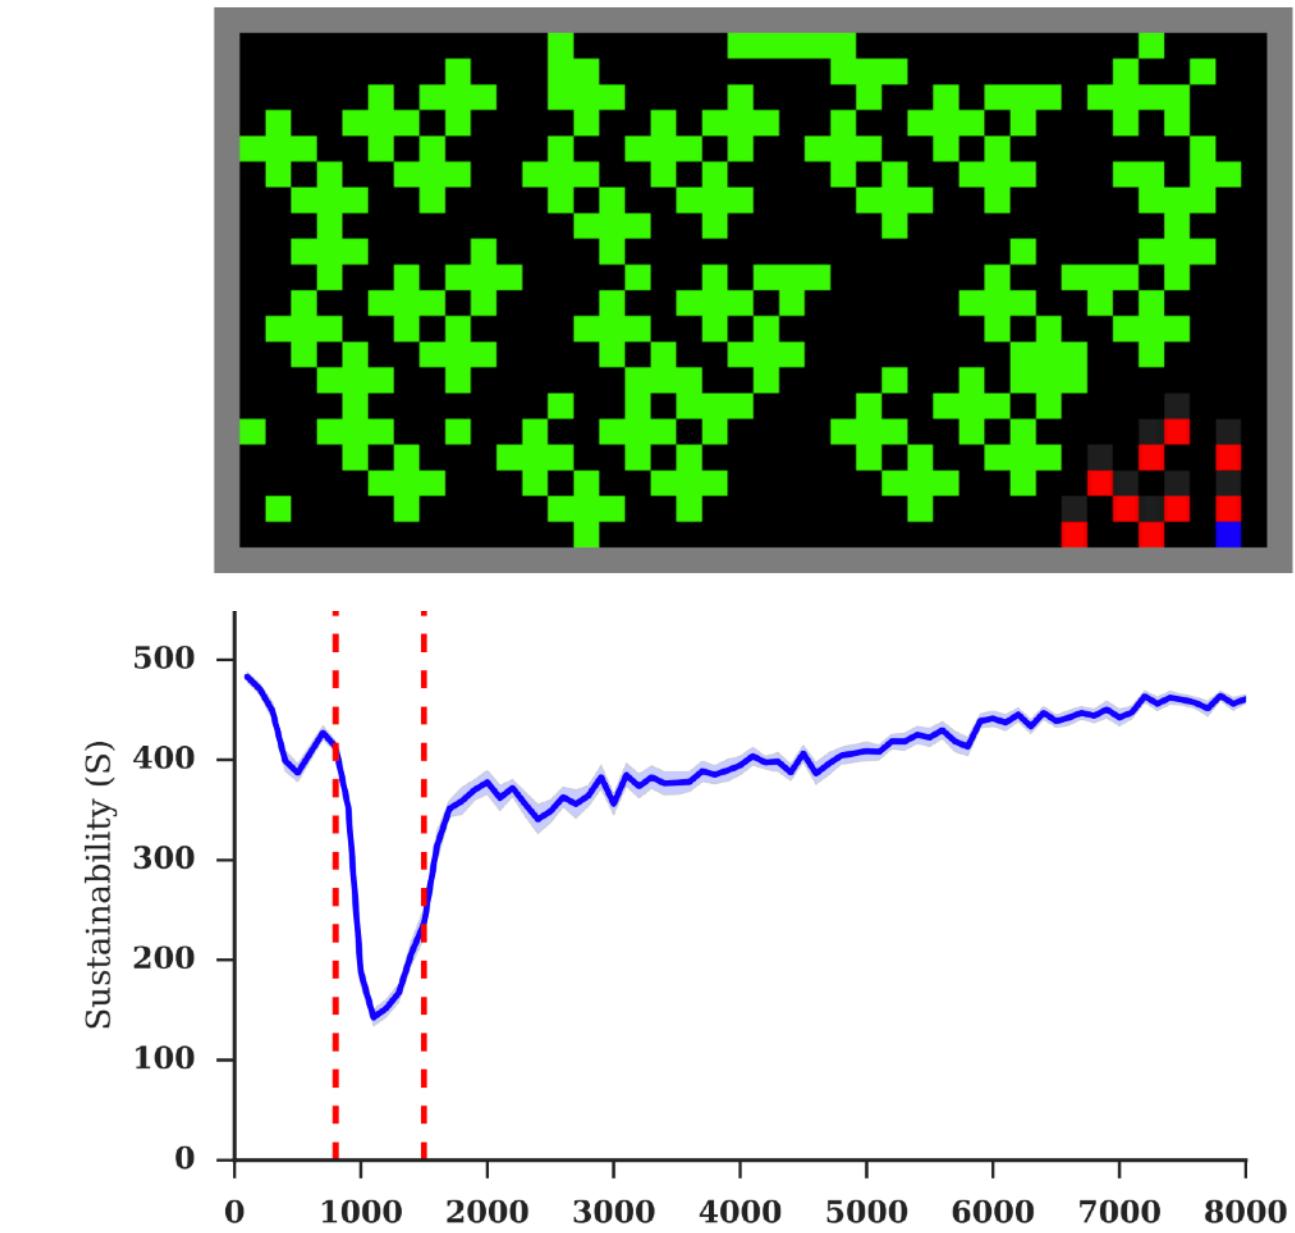
Barfuss et al., 2020



→ a semi-formal method

Deep MARL for common-pool resource appropriation

Perolat et al, 2017



Why learning dynamics? → Applications

Collective learning dynamics for solving MARL challenges

Why learning dynamics? → Applications

Collective learning dynamics for solving MARL challenges

MARL Challenges

There has been a lot of great work on multi-agent reinforcement learning (MARL) in the past decade, but significant challenges remain, including:

- the difficulty of learning an optimal model/policy from a partial signal,
- learning to cooperate/compete in non-stationary environments with distributed, simultaneously learning agents,
- the interplay between abstraction and influence of other agents,
- the exploration vs. exploitation dilemma,
- the scalability and effectiveness of learning algorithms,
- avoiding social dilemmas, and
- learning emergent communication.

[COMARL Symposium Website](#)

Why learning dynamics? → Applications

Collective learning dynamics for solving MARL challenges

MARL Challenges

There has been a lot of great work on multi-agent reinforcement learning (MARL) in the past decade, but significant challenges remain, including:

- the difficulty of learning an optimal model/policy from a partial signal,
- learning to cooperate/compete in non-stationary environments with distributed, simultaneously learning agents,
- the interplay between abstraction and influence of other agents,
- the exploration vs. exploitation dilemma,
- the scalability and effectiveness of learning algorithms,
- avoiding social dilemmas, and
- learning emergent communication.

[COMARL Symposium Website](#)

Application domains

- Interpretability
of multi-agent learning systems
- Analysis of strategic interactions
of multiple learning agents in changing environments
- Hyper-parameter tuning
- Design of new algorithms

Why learning dynamics? → Applications

Collective learning dynamics for solving MARL challenges

MARL Challenges

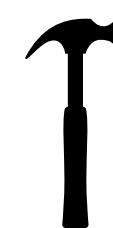
There has been a lot of great work on multi-agent reinforcement learning (MARL) in the past decade, but significant challenges remain, including:

- the difficulty of learning an optimal model/policy from a partial signal,
- learning to cooperate/compete in non-stationary environments with distributed, simultaneously learning agents,
- the interplay between abstraction and influence of other agents,
- the exploration vs. exploitation dilemma,
- the scalability and effectiveness of learning algorithms,
- avoiding social dilemmas, and
- learning emergent communication.

Application domains

- Interpretability
of multi-agent learning systems
- Analysis of strategic interactions
of multiple learning agents in changing environments
- Hyper-parameter tuning
- Design of new algorithms

[COMARL Symposium Website](#)



Problem-Method Fit ←
How exactly?

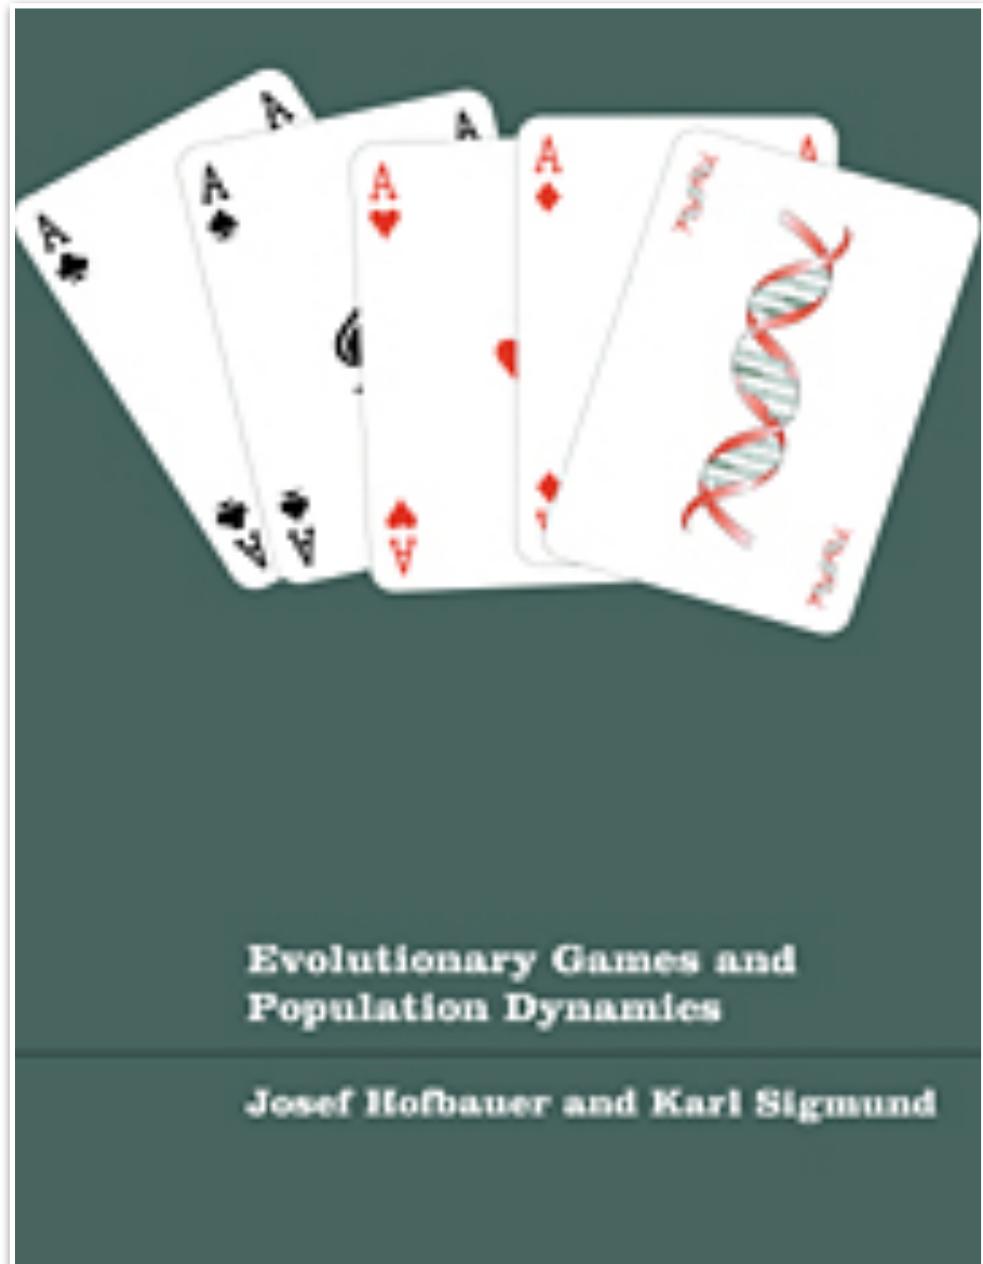
The challenge

Disciplinary divides a clustering attempt

The challenge

Disciplinary divides a clustering attempt

Biology
**Evolutionary Game
Theory**

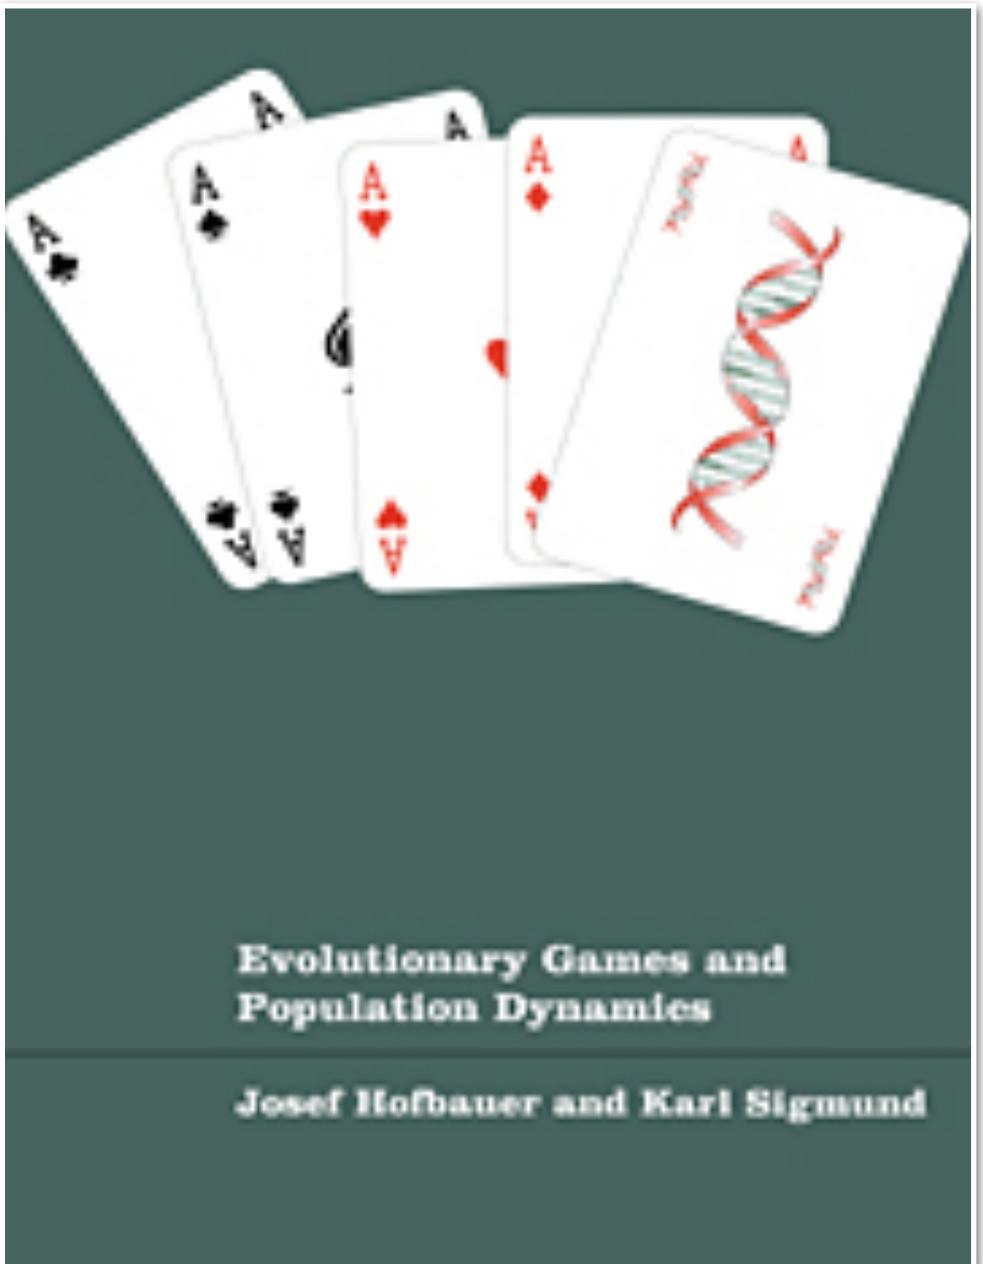


Hofbauer & Sigmund, 1998

The challenge

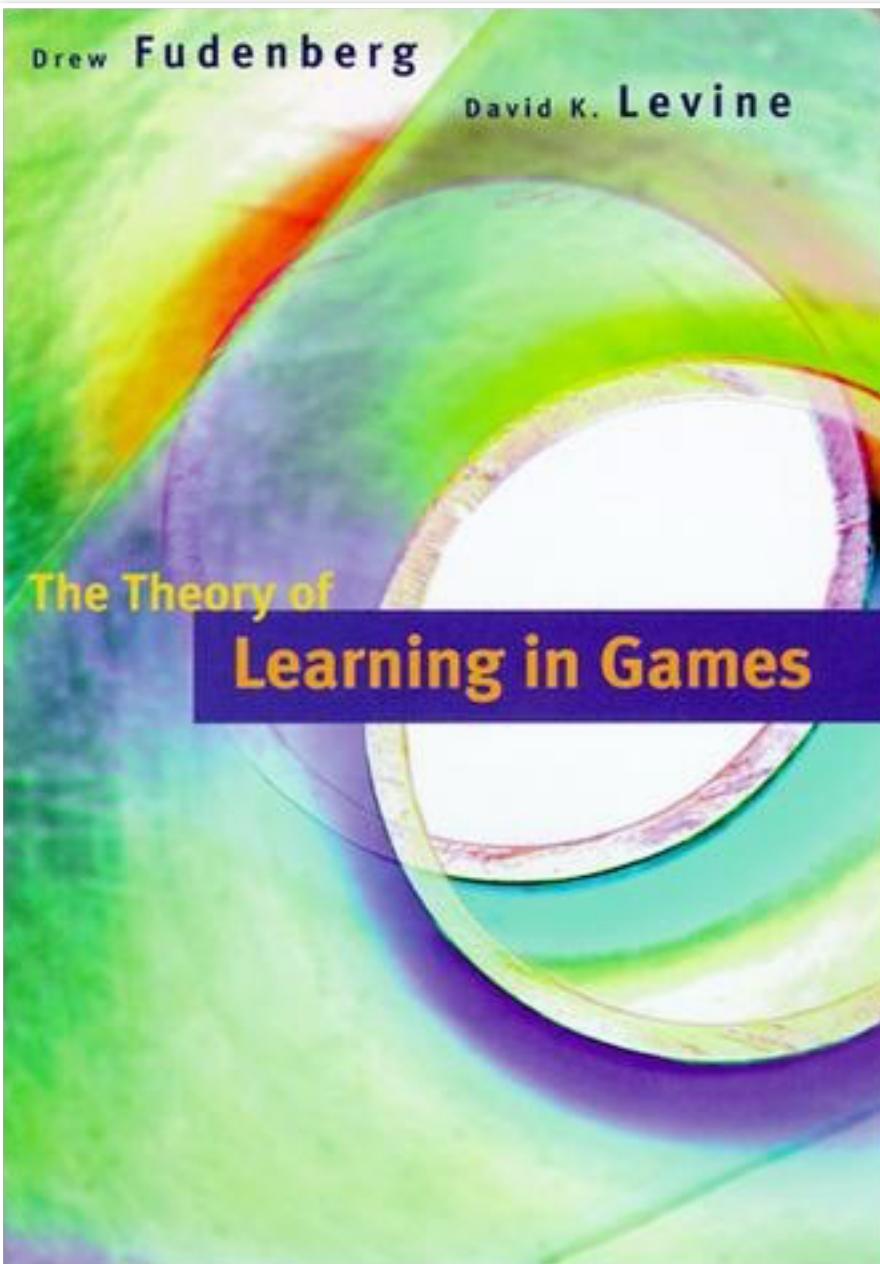
Disciplinary divides a clustering attempt

Biology
Evolutionary Game Theory



Hofbauer & Sigmund, 1998

Economics
Learning in Games

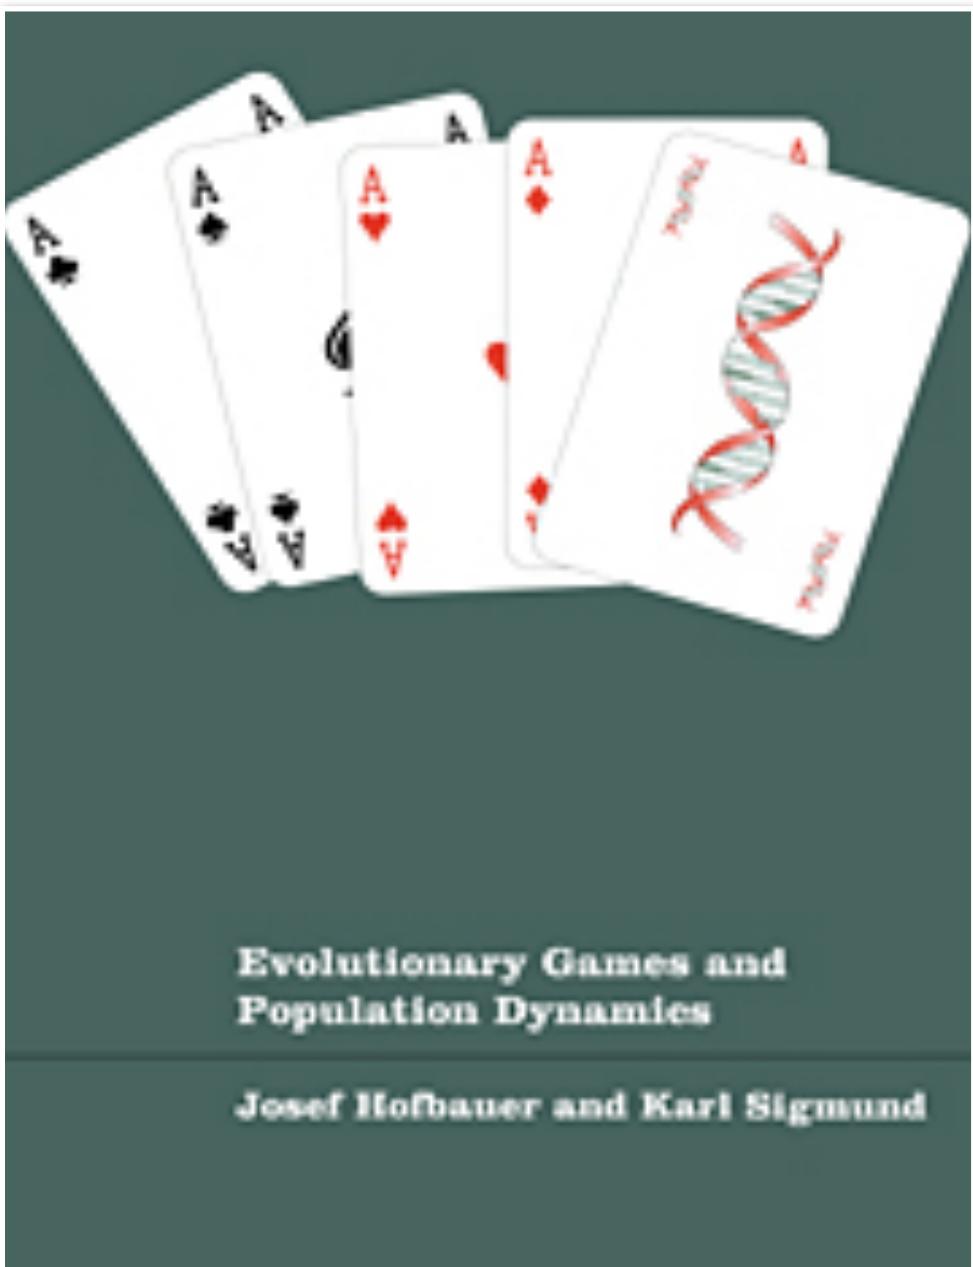


Fudenberg & Levine, 1998

The challenge

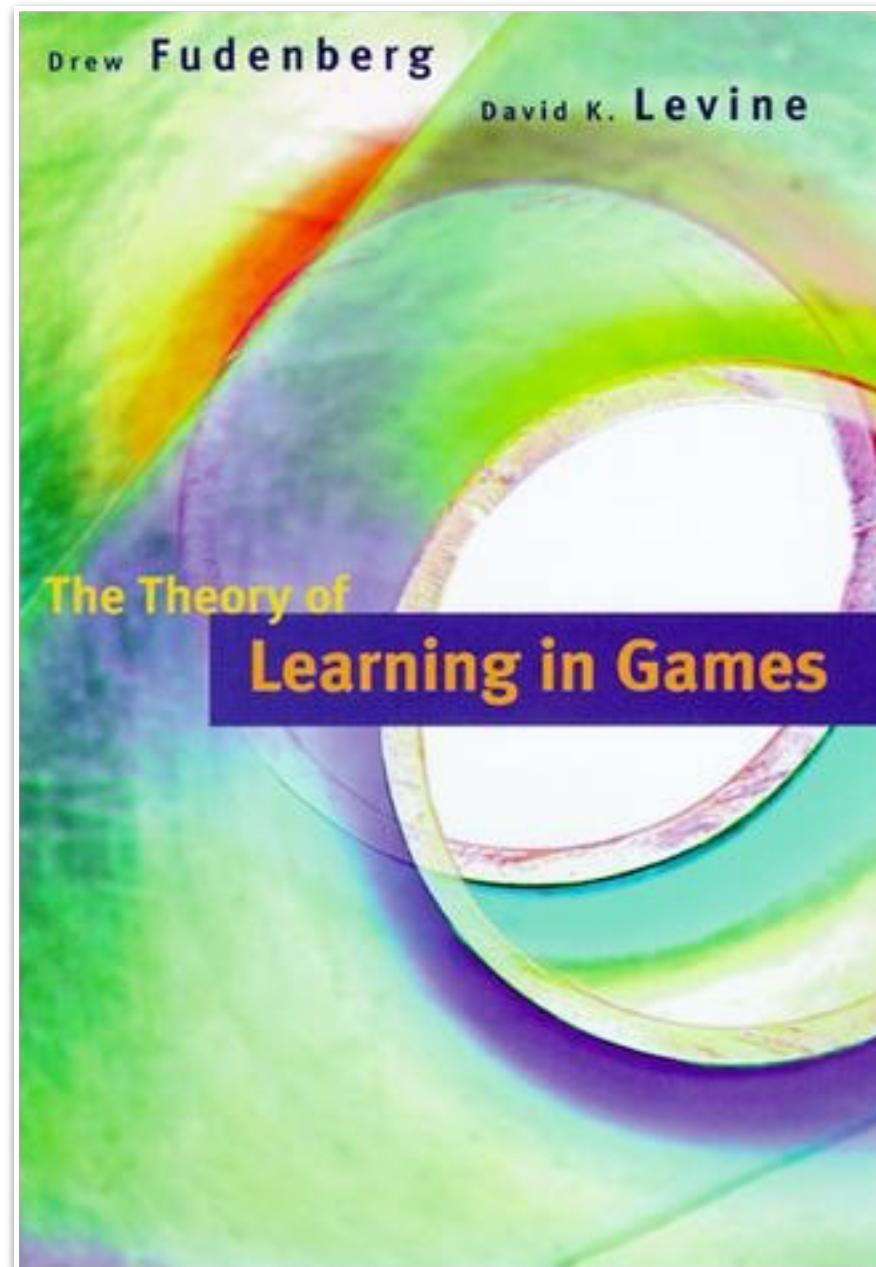
Disciplinary divides a clustering attempt

Biology
Evolutionary Game Theory



Hofbauer & Sigmund, 1998

Economics
Learning in Games



Fudenberg & Levine, 1998

Computer Science
Machine Learning Dynamics

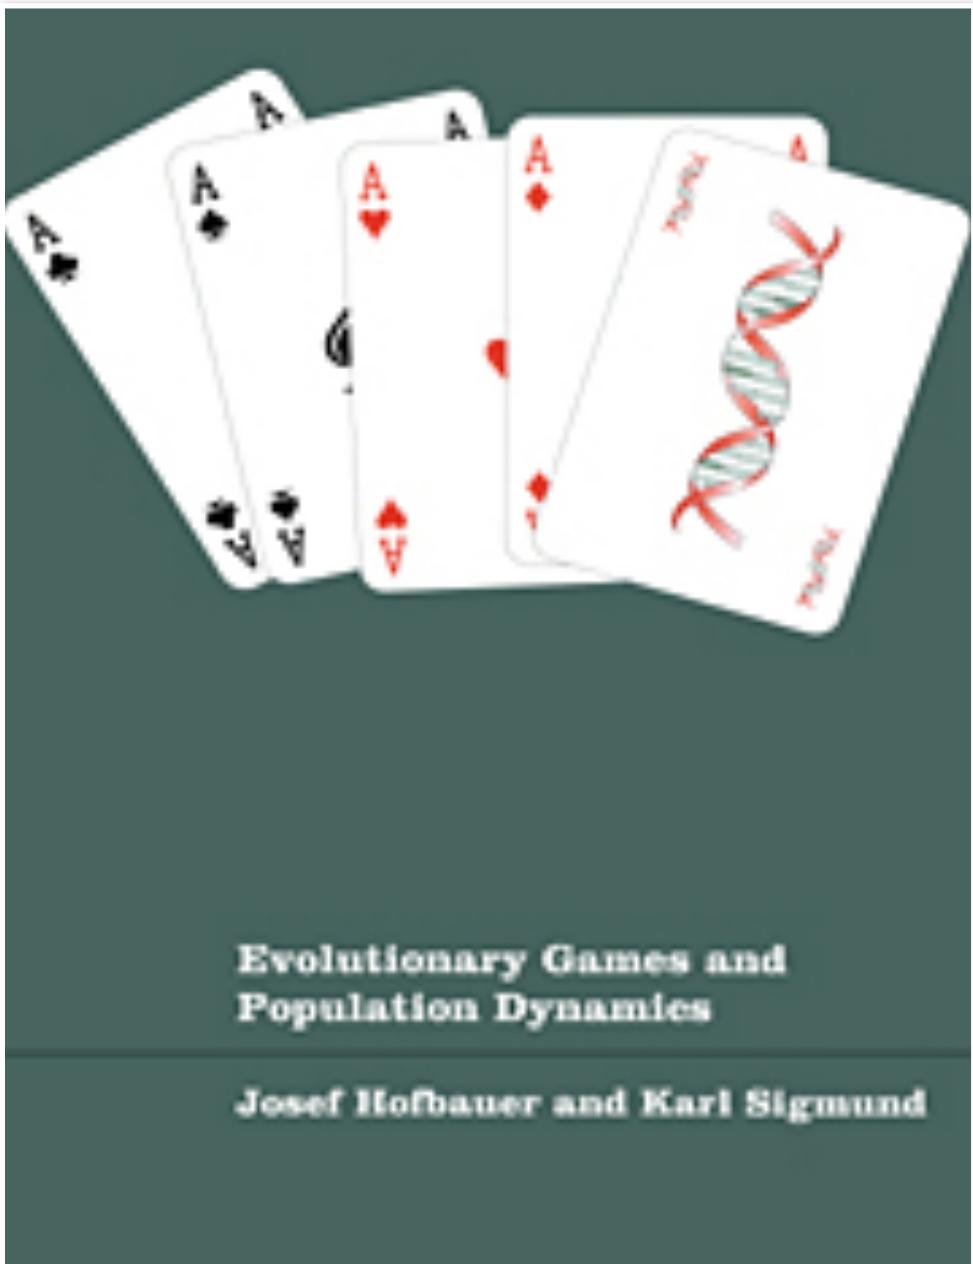


Bloembergen et al., 2015

The challenge

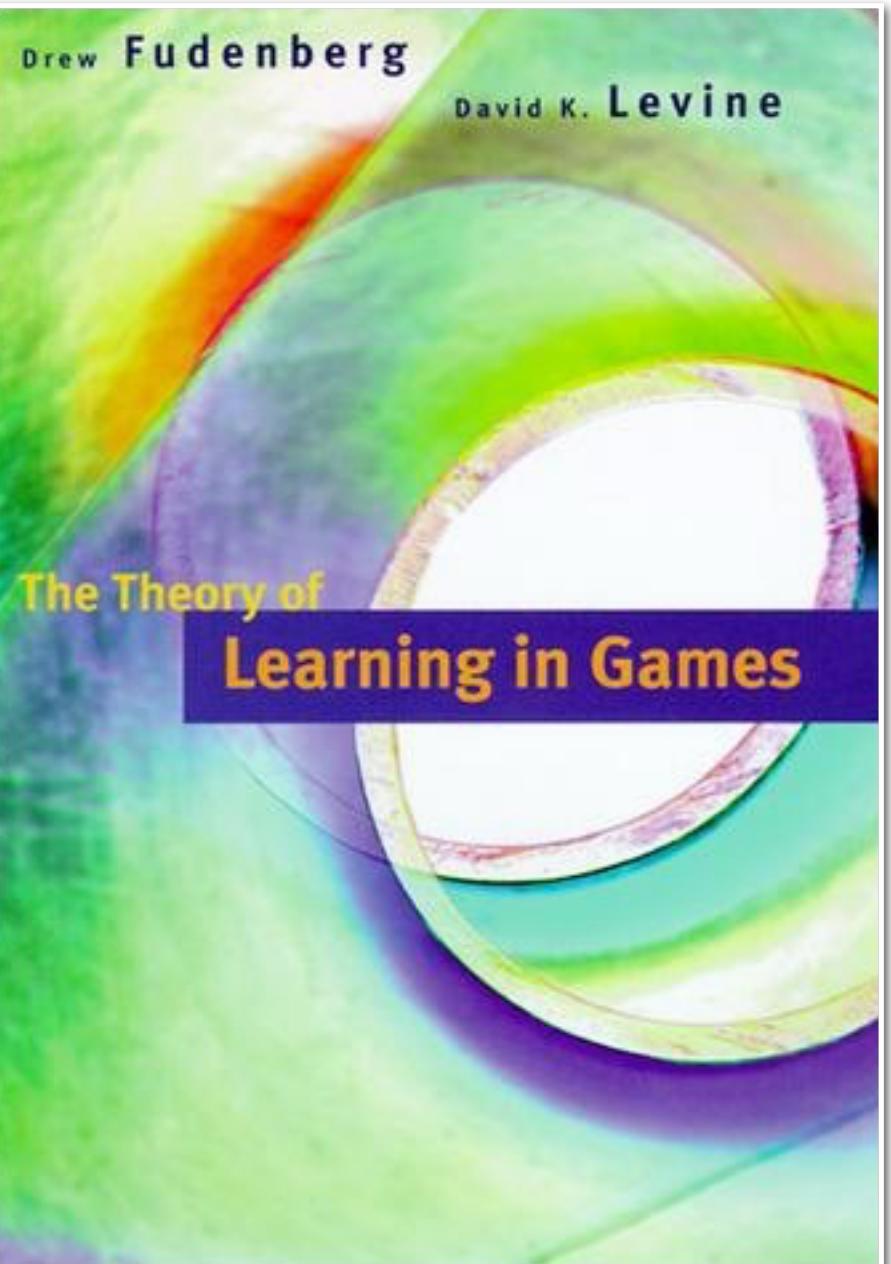
Disciplinary divides a clustering attempt

Biology Evolutionary Game Theory



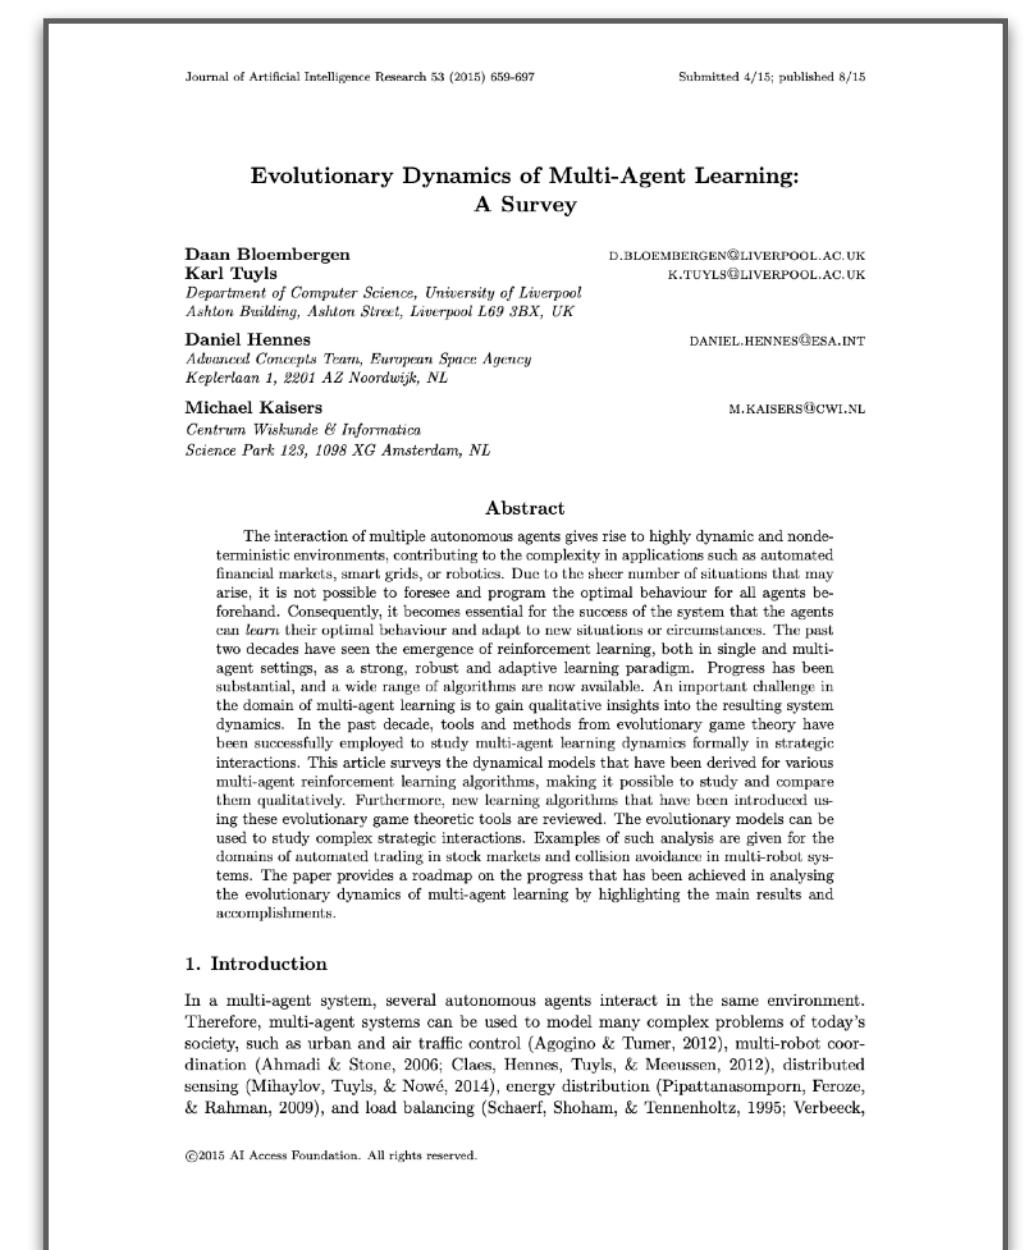
Hofbauer & Sigmund, 1998

Economics Learning in Games



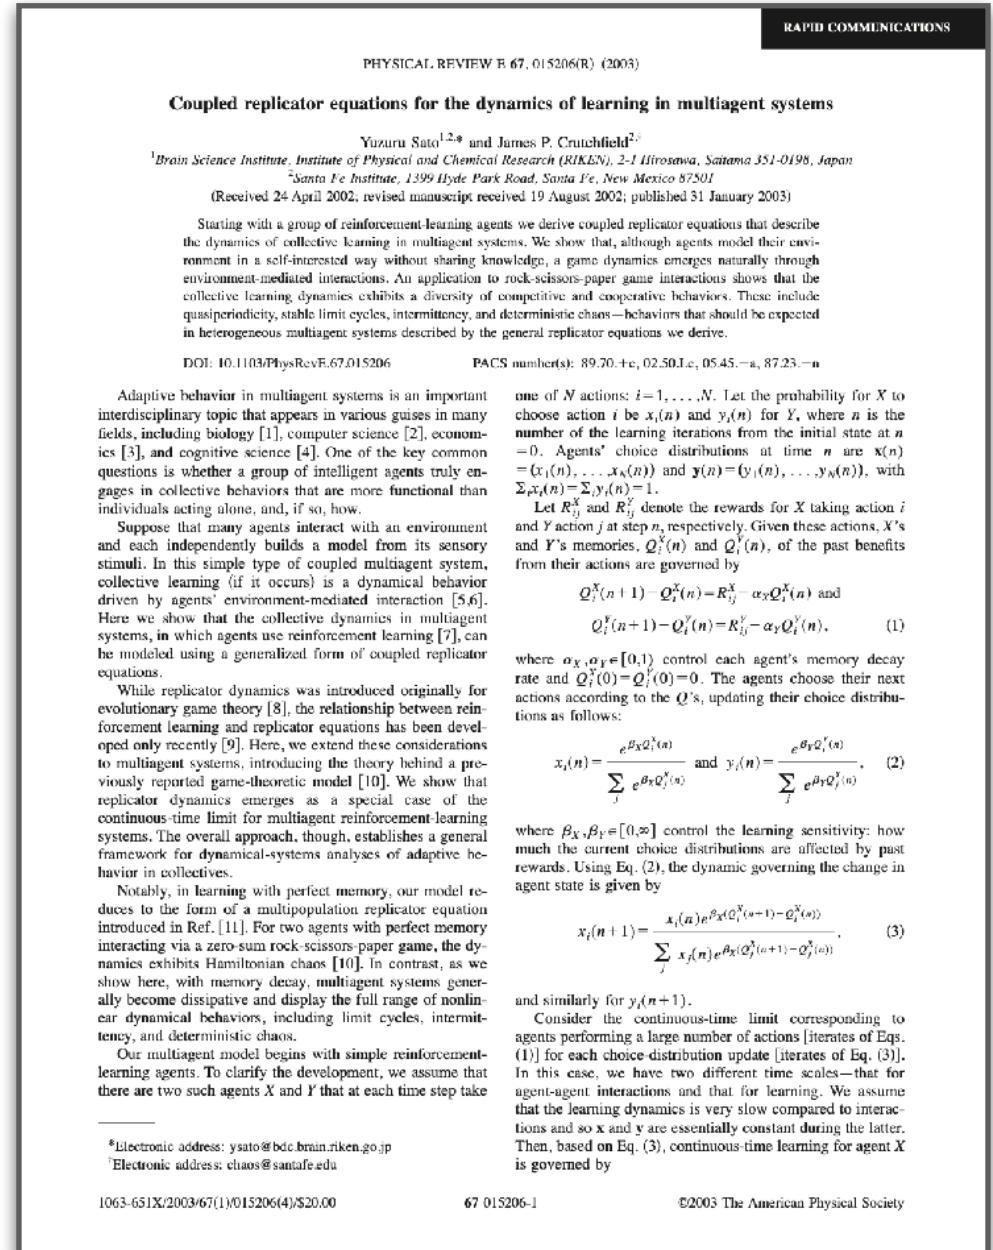
Fudenberg & Levine, 1998

Computer Science Machine Learning Dynamics



Bloembergen et al., 2015

Physics Statistical mechanics learning

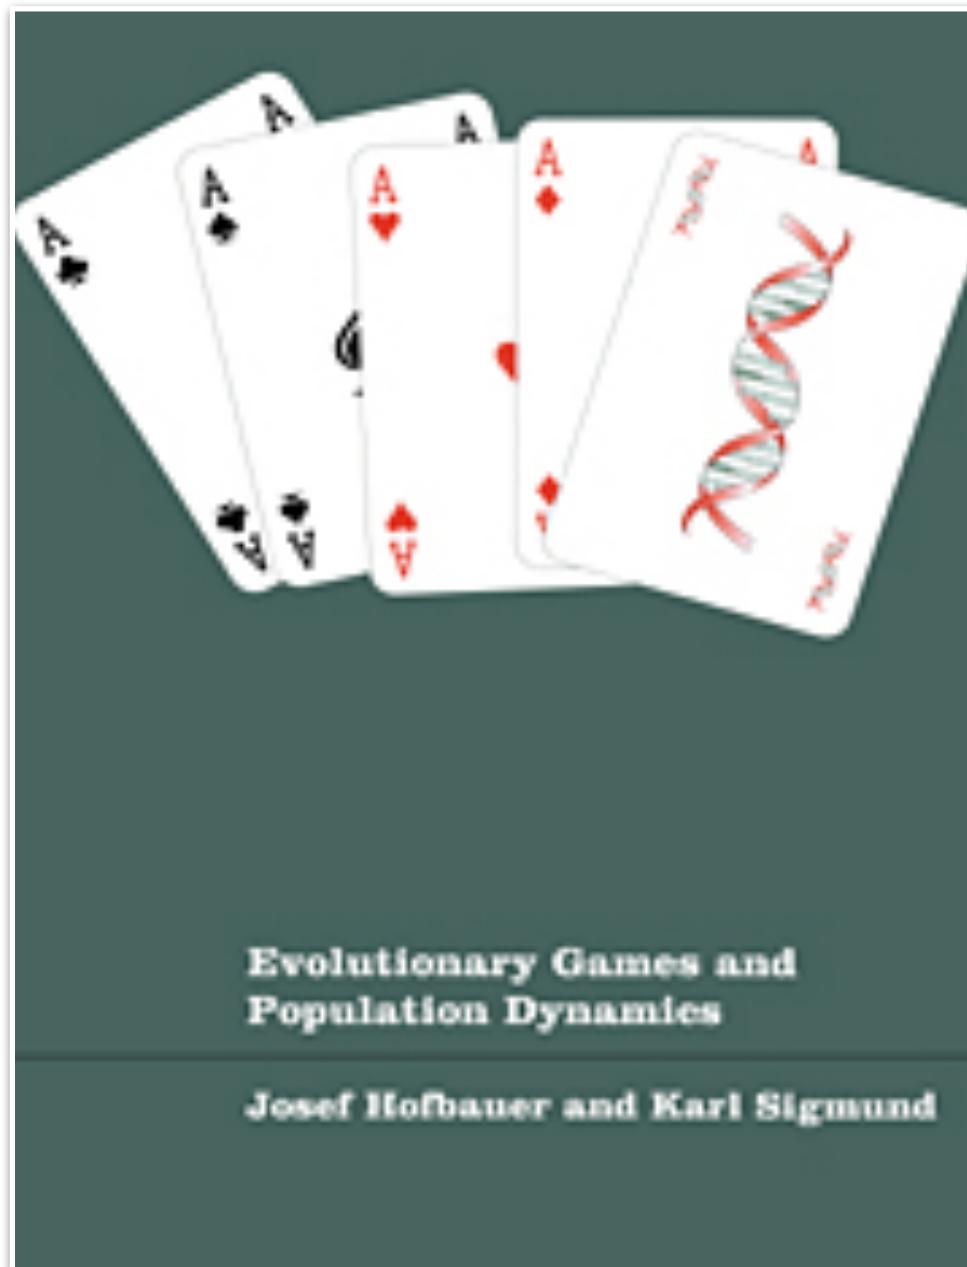


Sato & Crutchfield, 2003

The challenge

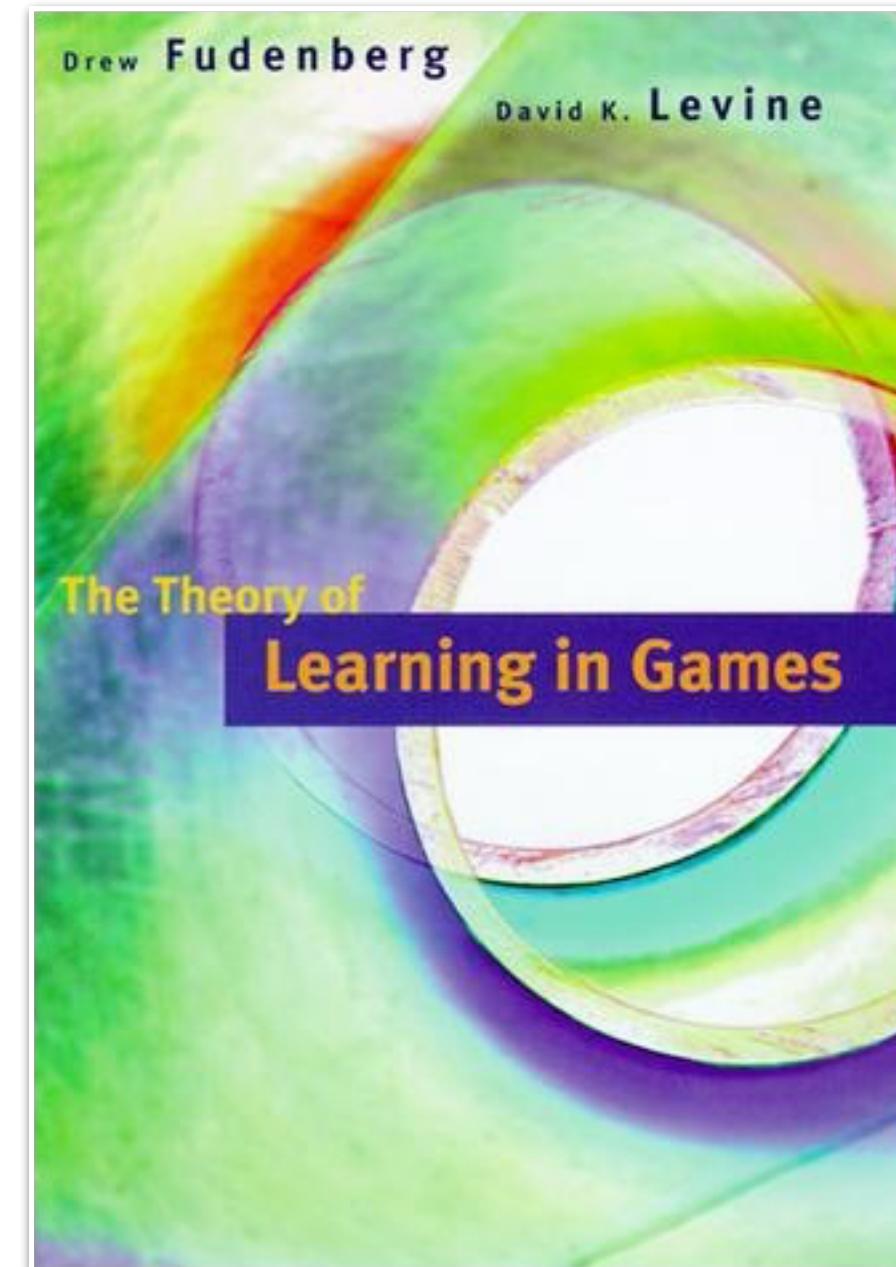
Disciplinary divides a clustering attempt

Biology Evolutionary Game Theory



Hofbauer & Sigmund, 1998

Economics Learning in Games



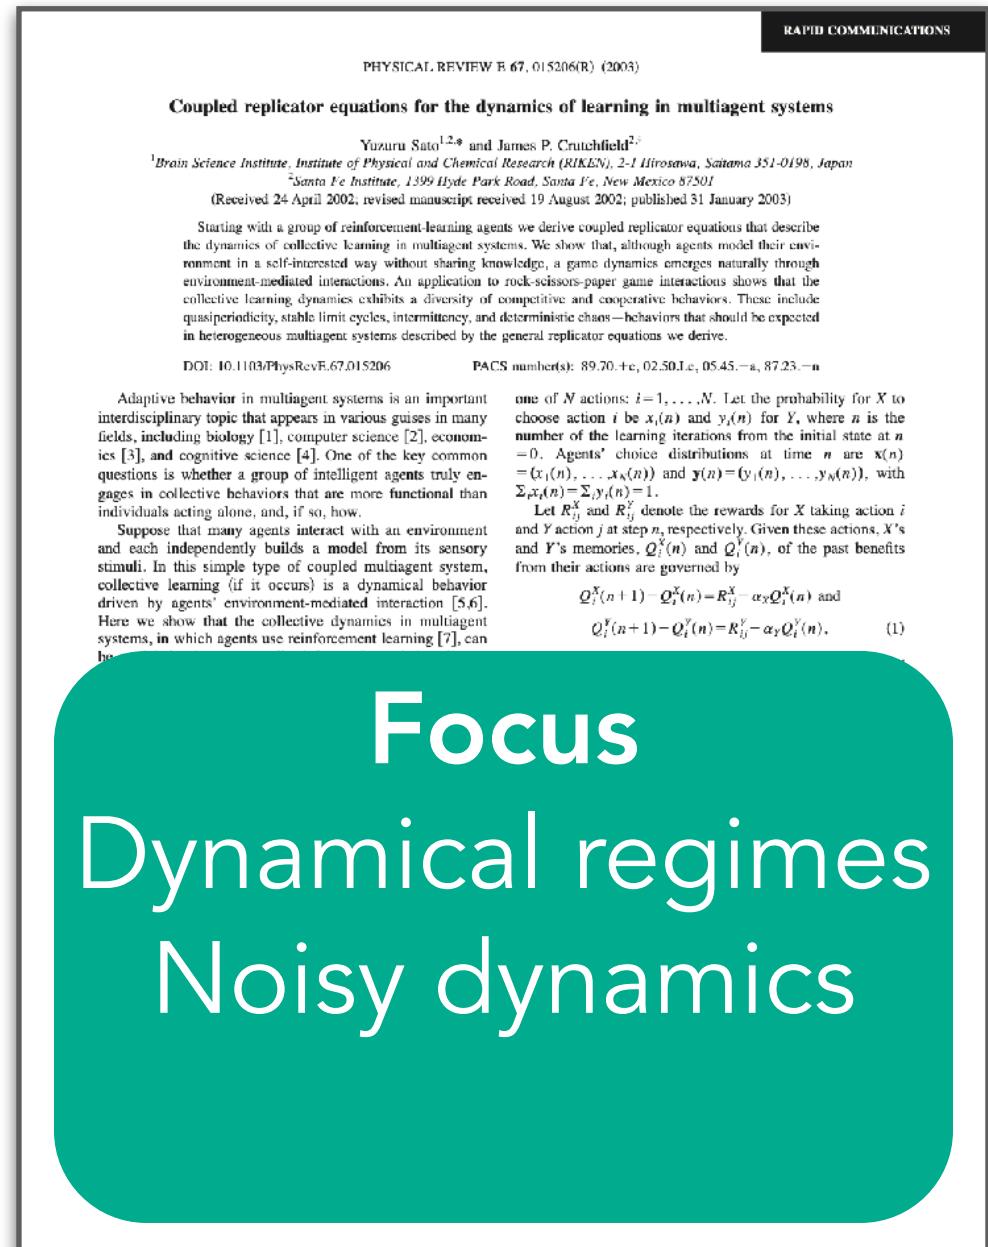
Fudenberg & Levine, 1998

Computer Science Machine Learning Dynamics



Bloembergen et al., 2015

Physics Statistical mechanics learning



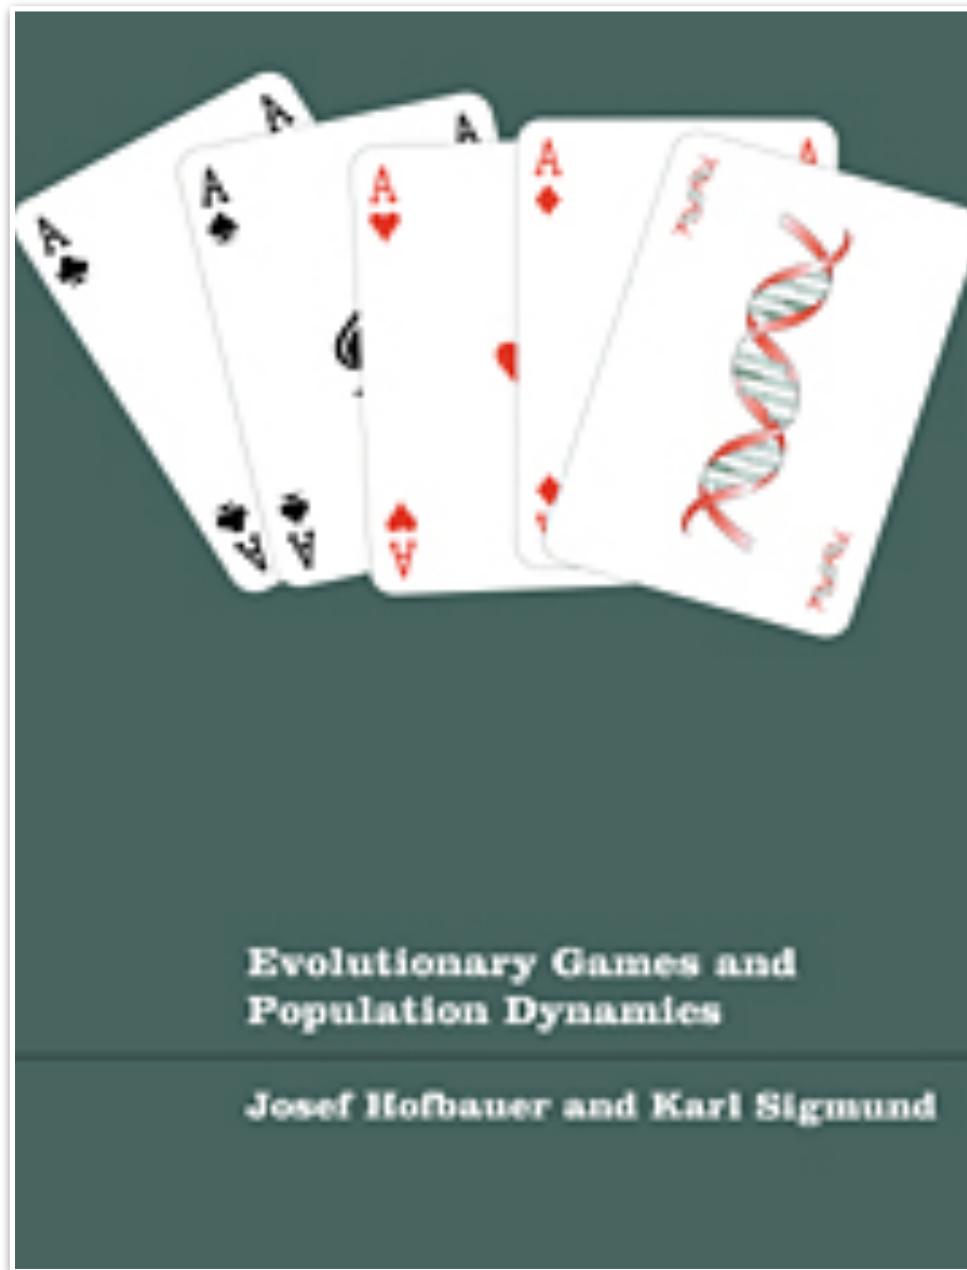
Sato & Crutchfield, 2003

Focus
Dynamical regimes
Noisy dynamics

The challenge

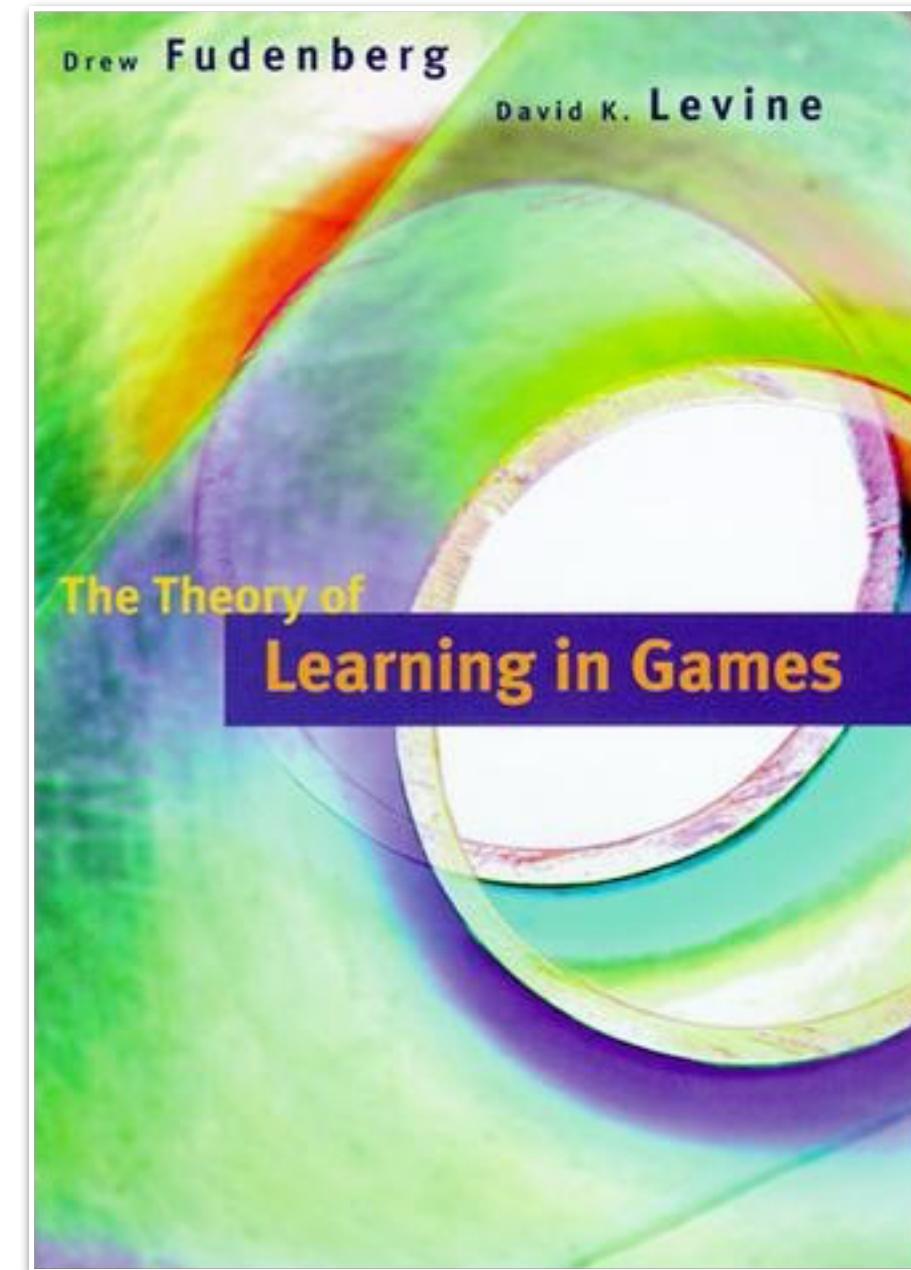
Disciplinary divides a clustering attempt

Biology Evolutionary Game Theory



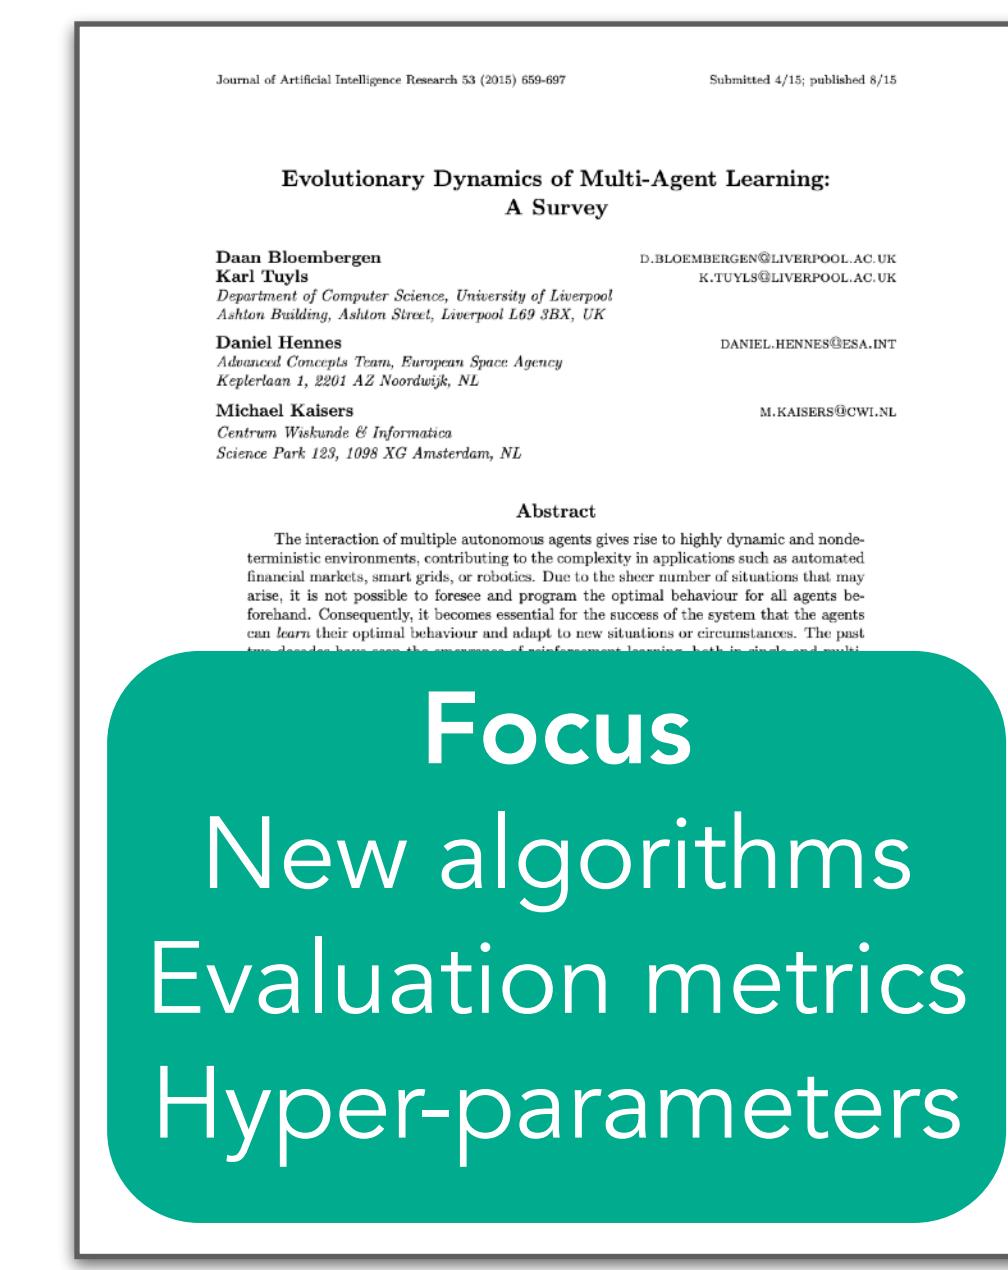
Hofbauer & Sigmund, 1998

Economics Learning in Games



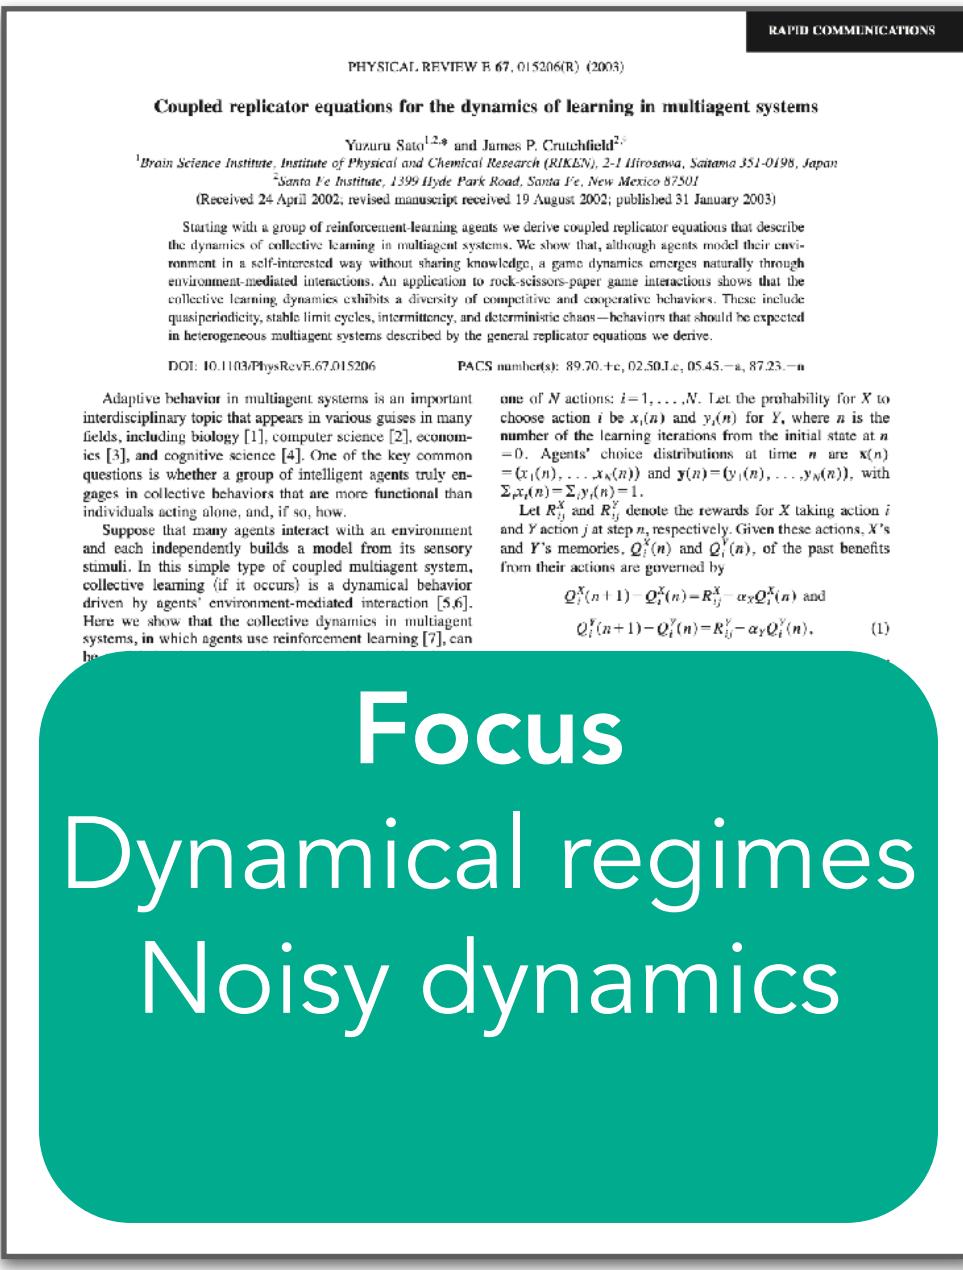
Fudenberg & Levine, 1998

Computer Science Machine Learning Dynamics



Bloembergen et al., 2015

Physics Statistical mechanics learning

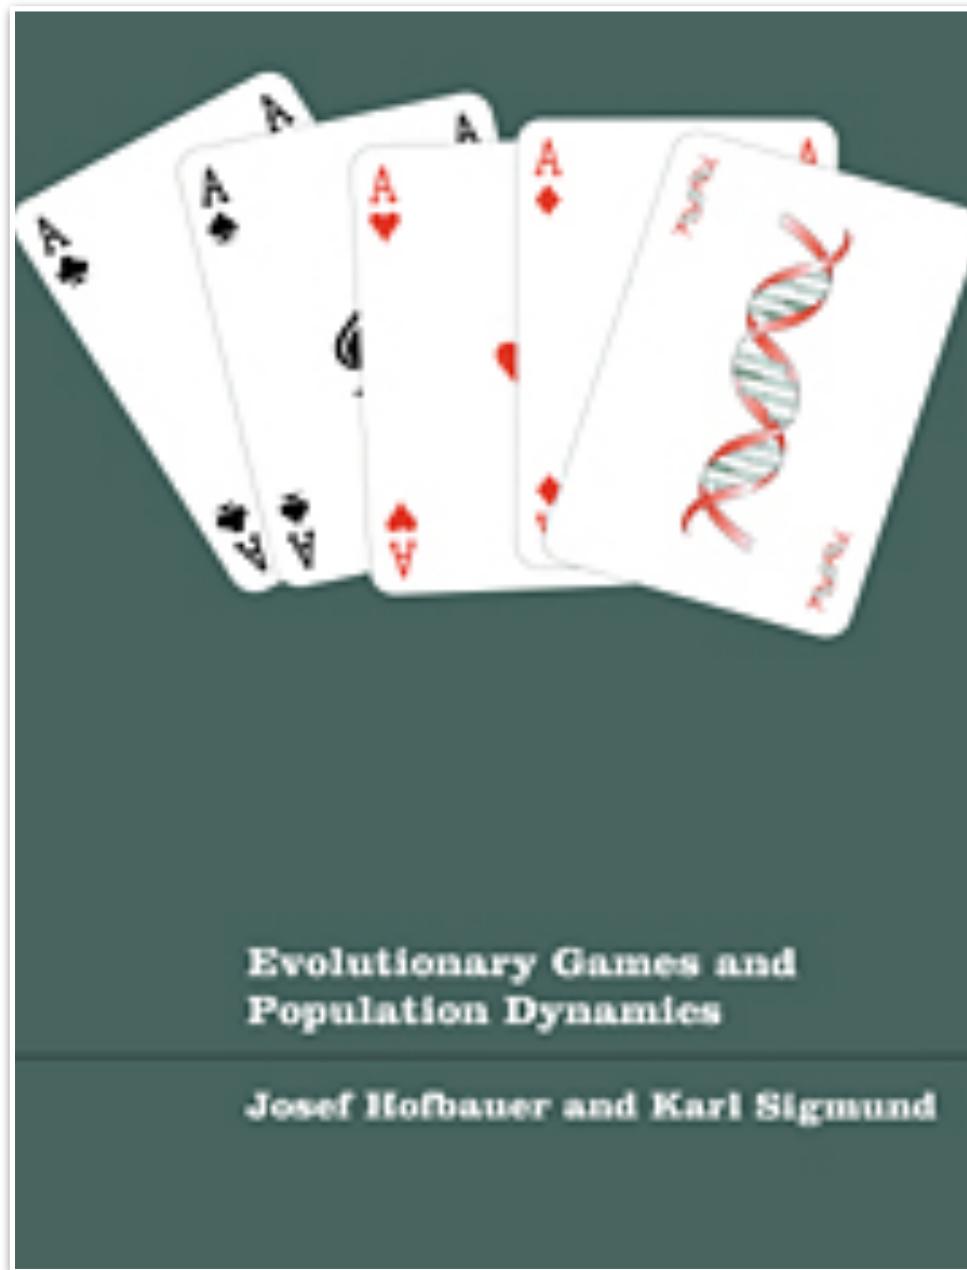


Sato & Crutchfield, 2003

The challenge

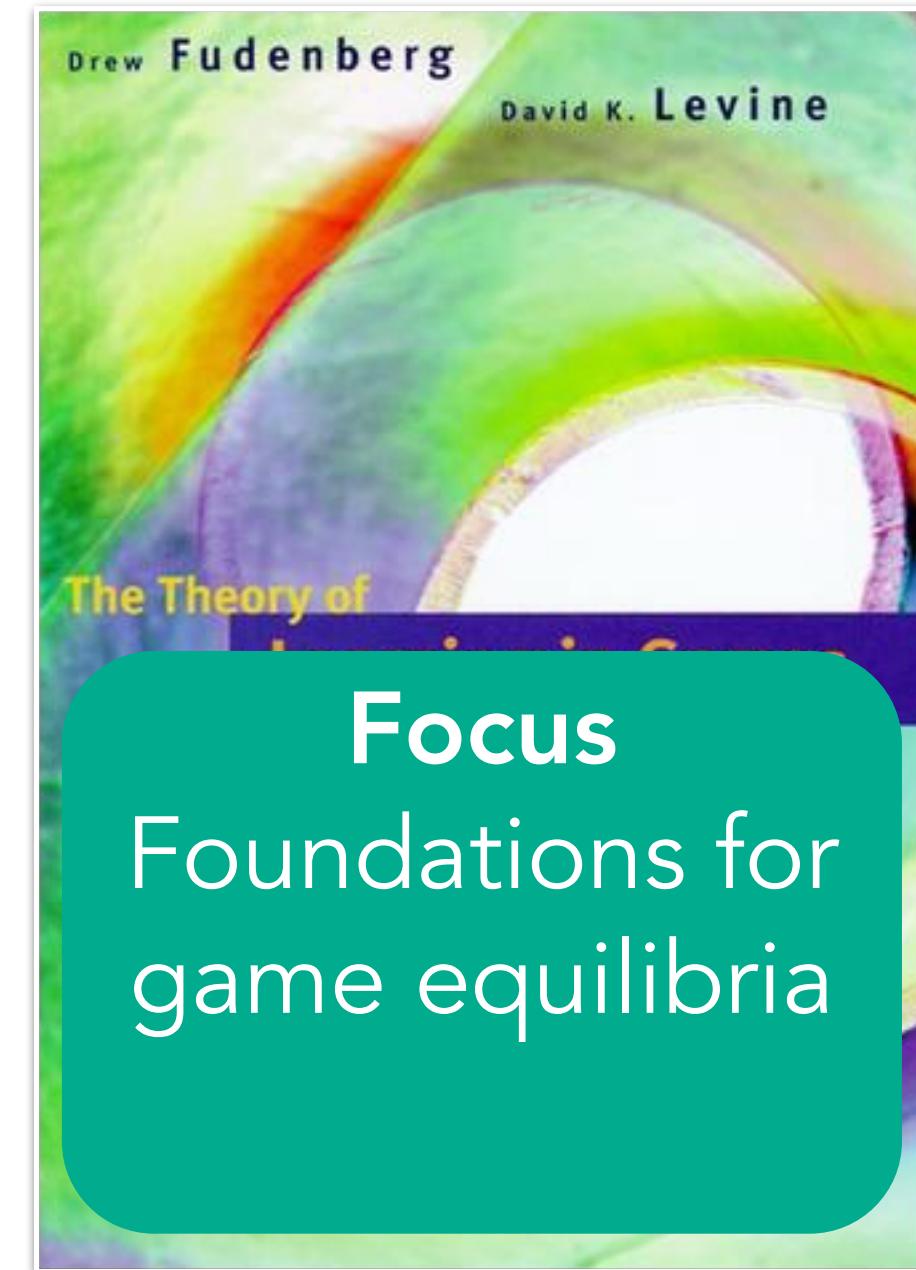
Disciplinary divides a clustering attempt

Biology Evolutionary Game Theory



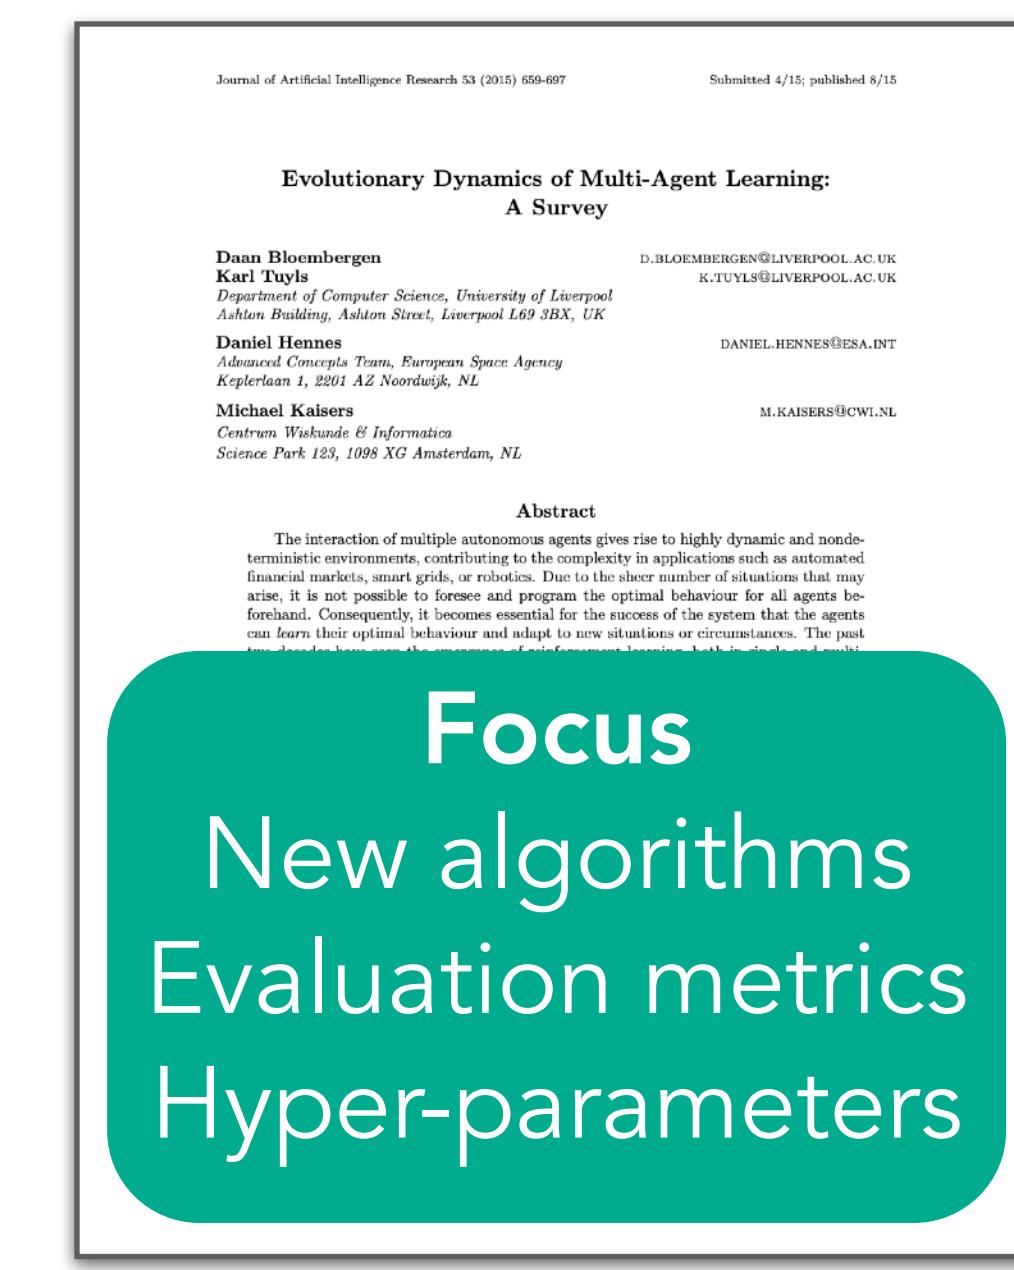
Hofbauer & Sigmund, 1998

Economics Learning in Games



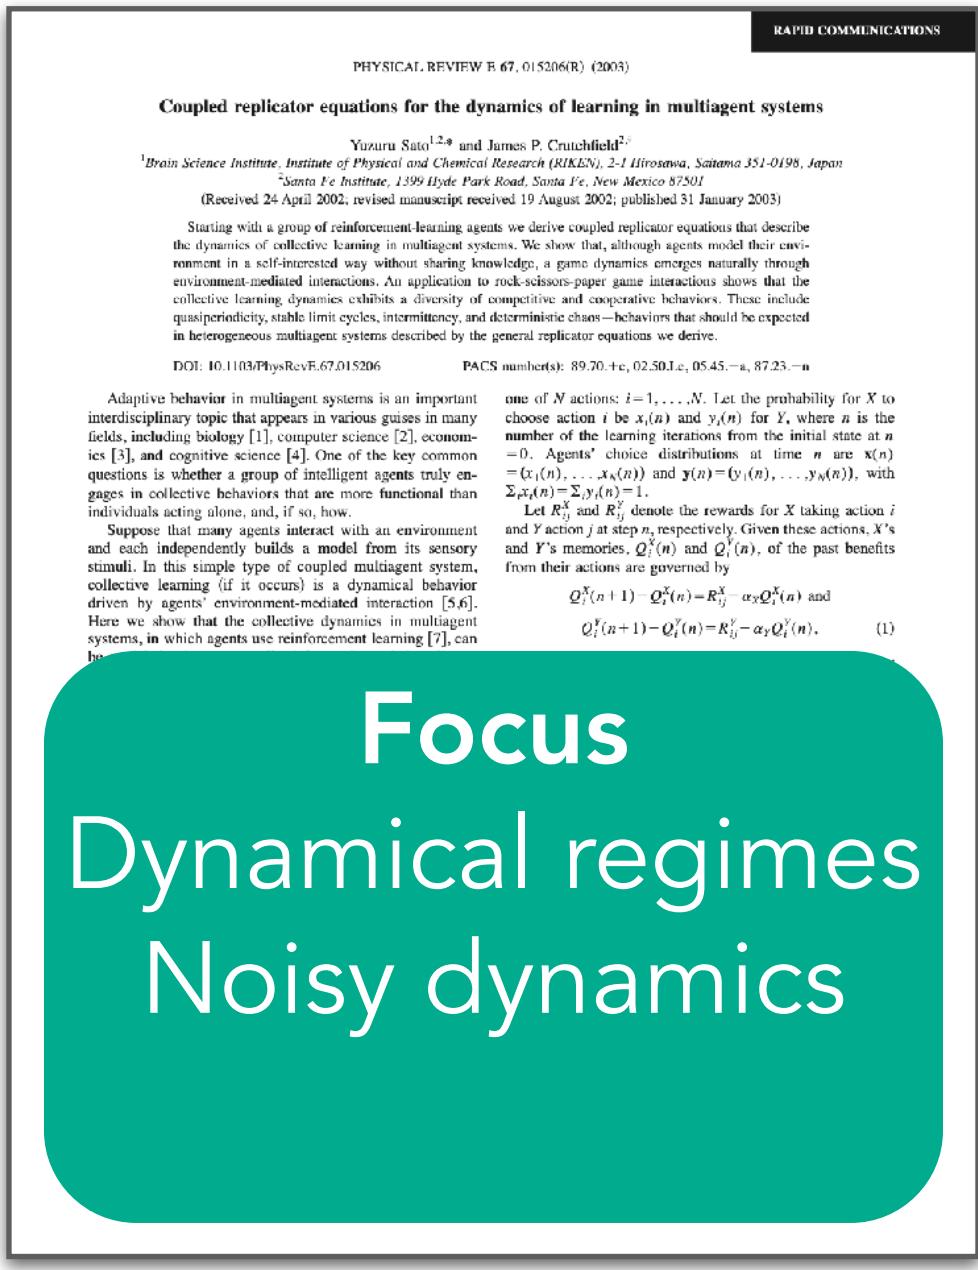
Fudenberg & Levine, 1998

Computer Science Machine Learning Dynamics



Bloembergen et al., 2015

Physics Statistical mechanics learning



Sato & Crutchfield, 2003

The challenge

Disciplinary divides a clustering attempt

Biology Evolutionary Game Theory



Focus
Understanding biological evolution

Hofbauer & Sigmund, 1998

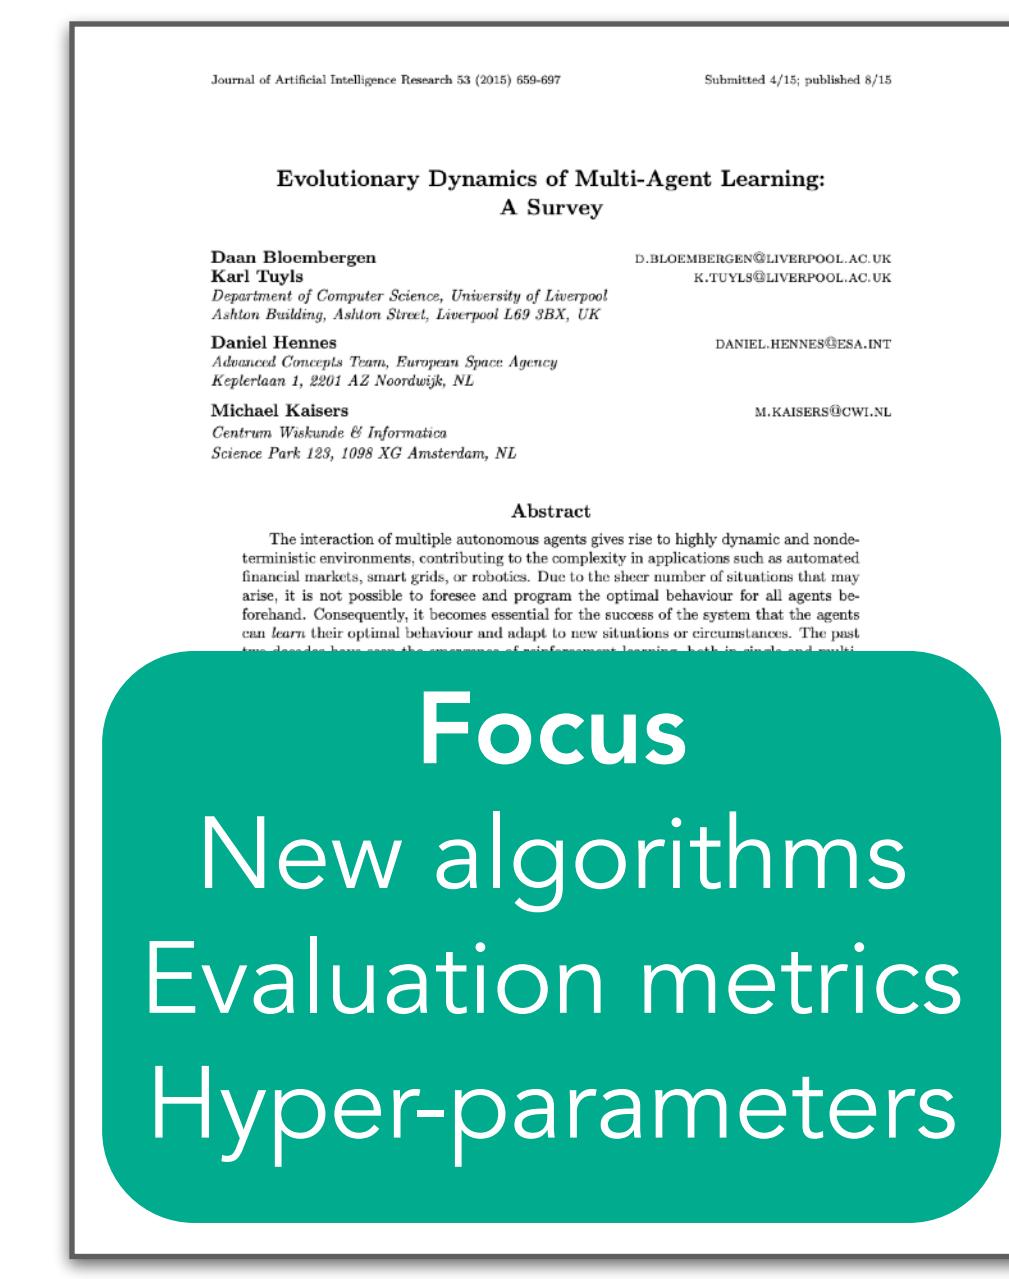
Economics Learning in Games



Focus
Foundations for game equilibria

Fudenberg & Levine, 1998

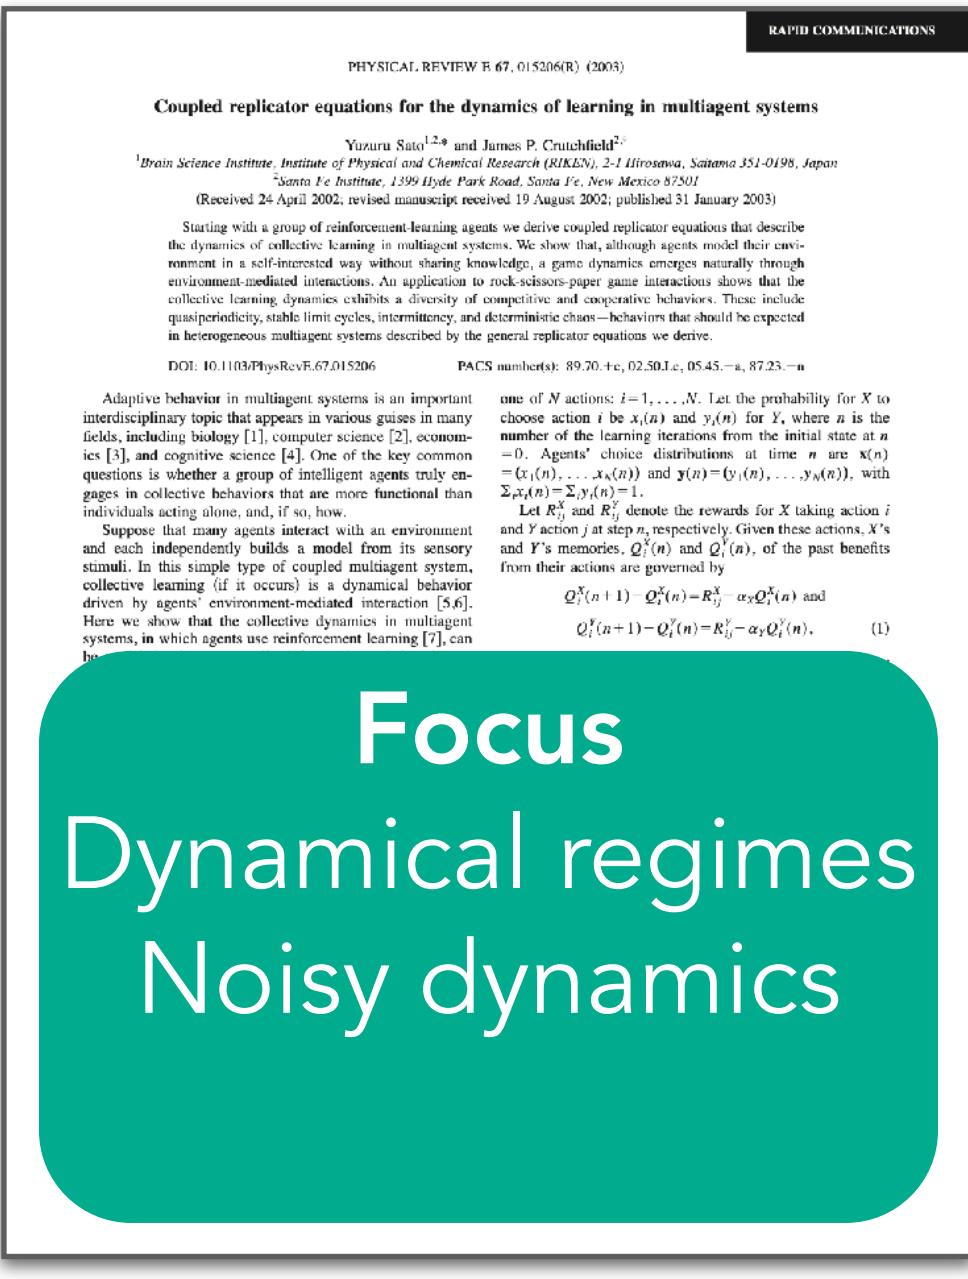
Computer Science Machine Learning Dynamics



Focus
New algorithms
Evaluation metrics
Hyper-parameters

Bloembergen et al., 2015

Physics Statistical mechanics learning



Focus
Dynamical regimes
Noisy dynamics

Sato & Crutchfield, 2003

The challenge

Disciplinary divides a clustering attempt

Biology Evolutionary Game Theory



Focus
Understanding biological evolution

Hofbauer & Sigmund, 1998

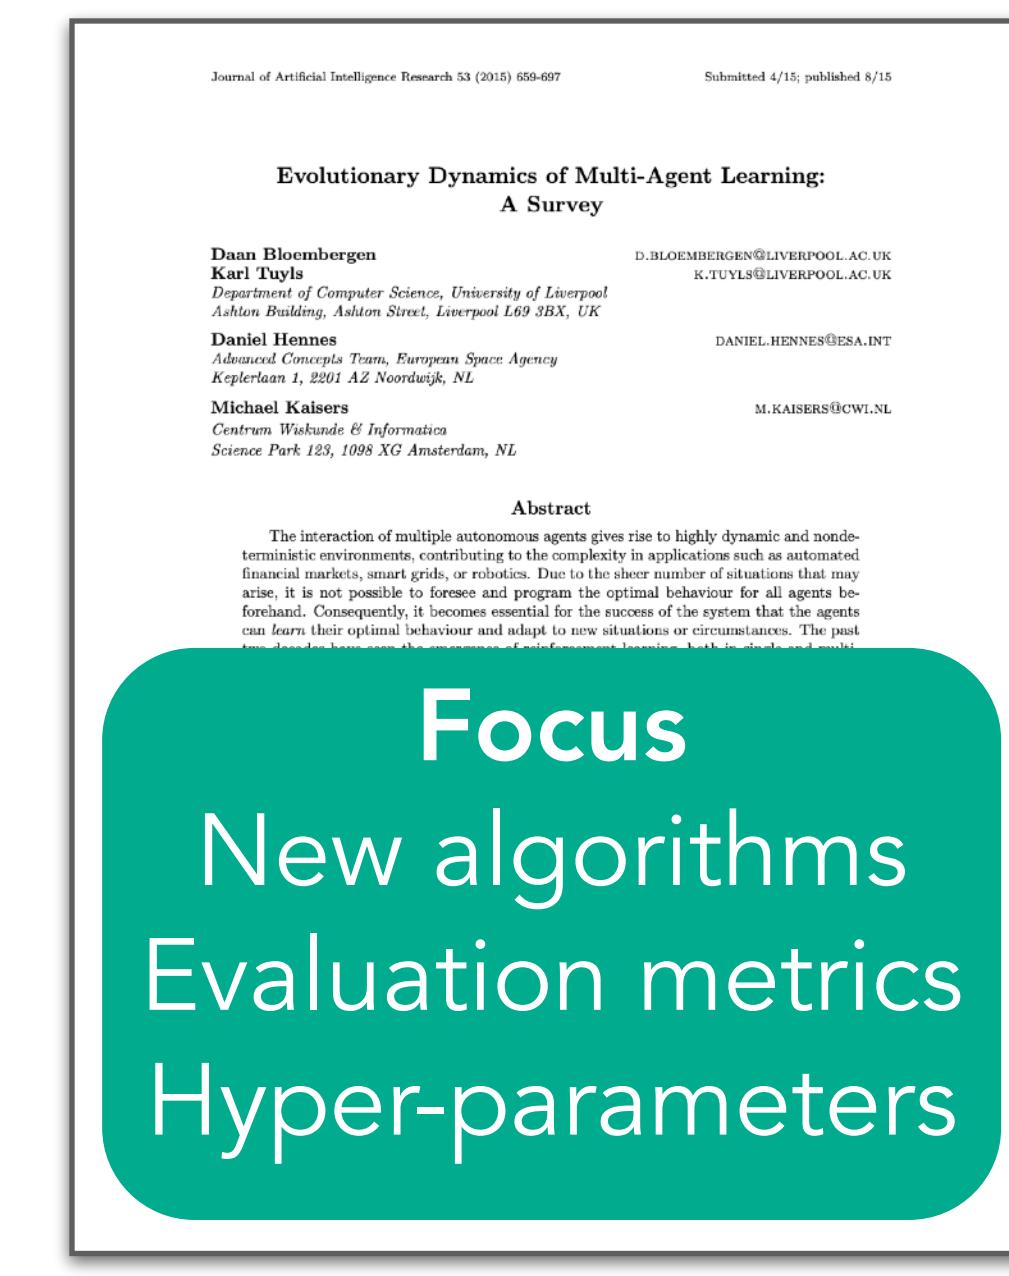
Economics Learning in Games



Focus
Foundations for game equilibria

Fudenberg & Levine, 1998

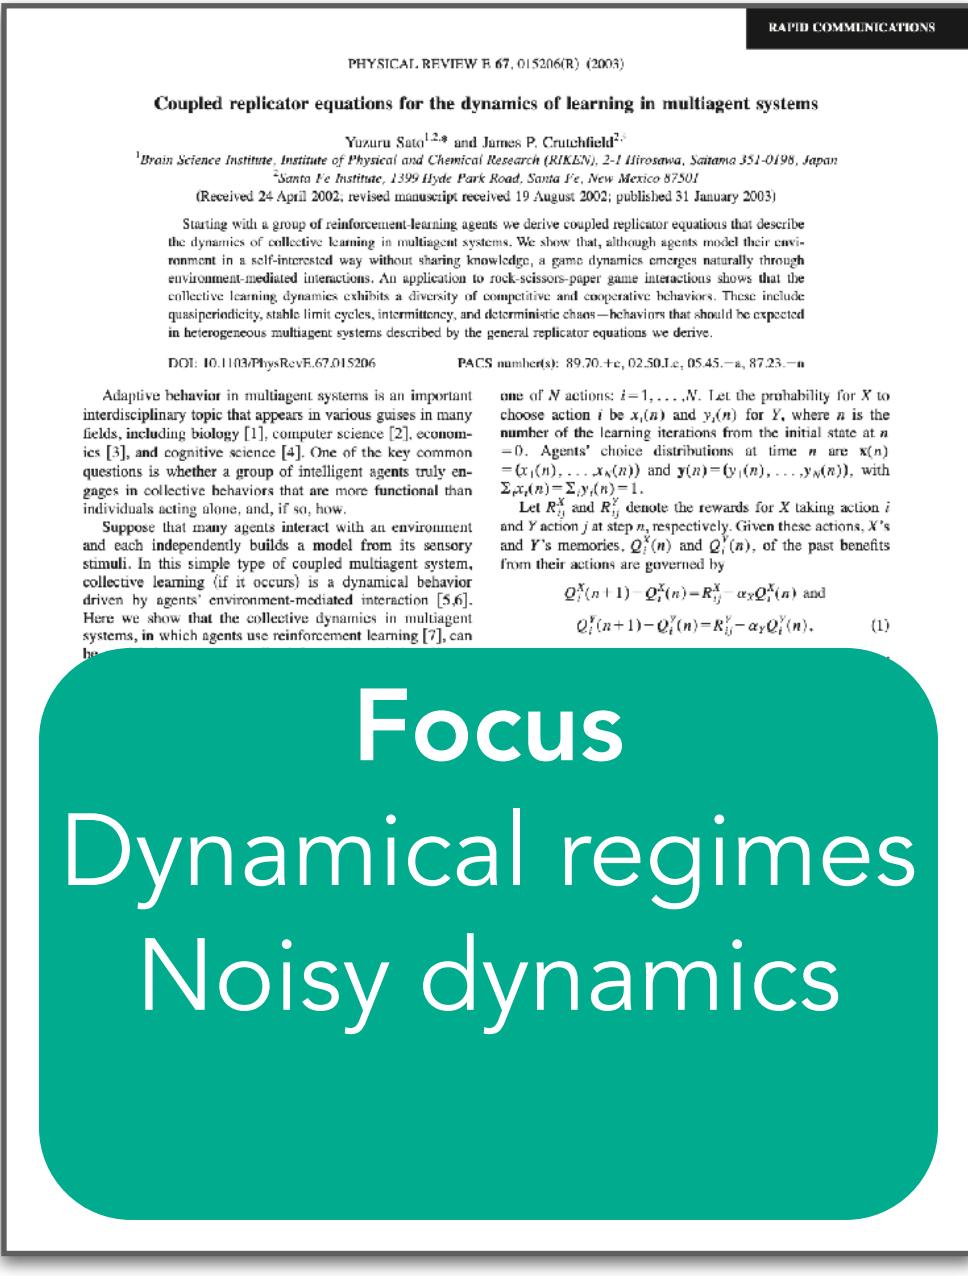
Computer Science Machine Learning Dynamics



Focus
New algorithms
Evaluation metrics
Hyper-parameters

Bloembergen et al., 2015

Physics Statistical mechanics learning



Focus
Dynamical regimes
Noisy dynamics

Sato & Crutchfield, 2003

→ Towards a unified treatment



Refining Environments

What do we want? ➔ a DynGym - a gym for learning *dynamics* ?

Refining Environments

What do we want? ➔ a DynGym - a gym for learning *dynamics* ?

- (Stylised) **single agent capabilities** (~Osband et al., 2019):
exploration, memory, credit assignment,...

Refining Environments

What do we want? ➔ a DynGym - a gym for learning *dynamics* ?

- (Stylised) **single agent capabilities** (~Osband et al., 2019):
exploration, memory, credit assignment,...
- (Stylised) **multi-agent capabilities**:
global exploration, coordination, cooperation (trust), competition (not exploitable),...

Refining Environments

What do we want? ➔ a DynGym - a gym for learning *dynamics* ?

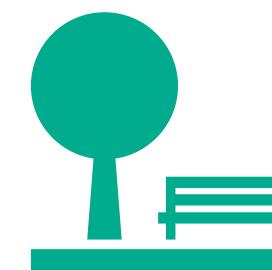
- (Stylised) **single agent capabilities** (~Osband et al., 2019):
exploration, memory, credit assignment,...
- (Stylised) **multi-agent capabilities**:
global exploration, coordination, cooperation (trust), competition (not exploitable),...
- (Stylised) **applications**:
AI safety (~Leike et al. 2017), AI governance (~The Anh Han et al. 2020),
Sustainability in social-ecological dilemmas (~Barfuss et al. 2020), ...

Refining Environments

What do we want? ➔ a DynGym - a gym for learning *dynamics* ?

- (Stylised) **single agent capabilities** (~Osband et al., 2019):
exploration, memory, credit assignment,...
- (Stylised) **multi-agent capabilities**:
global exploration, coordination, cooperation (trust), competition (not exploitable),...
- (Stylised) **applications**:
AI safety (~Leike et al. 2017), AI governance (~The Anh Han et al. 2020),
Sustainability in social-ecological dilemmas (~Barfuss et al. 2020), ...

➔ Benchmarks?



Refining Agents

Advancing agents / relaxing assumptions

Refining Agents

Advancing agents / relaxing assumptions

- RRL agents with **partial observability**

Refining Agents

Advancing agents / relaxing assumptions

- RRL agents with **partial observability**
- RRL agents with **stylised function approximations**

Refining Agents

Advancing agents / relaxing assumptions

- RRL agents with **partial observability**
- RRL agents with **stylised function approximations**
- RRL agents with **finite memory batch** > noisy dynamics

Refining Agents

Advancing agents / relaxing assumptions

- RRL agents with **partial observability**
- RRL agents with **stylised function approximations**
- RRL agents with **finite memory batch** > noisy dynamics
- RRL agents with **different learning paradigms** (imitation, belief, ...)

Refining Agents

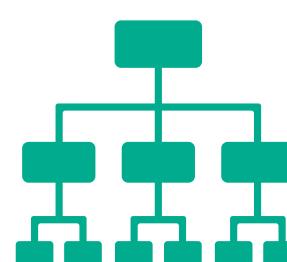
Advancing agents / relaxing assumptions

- RRL agents with **partial observability**
- RRL agents with **stylised function approximations**
- RRL agents with **finite memory batch** > noisy dynamics
- RRL agents with **different learning paradigms** (imitation, belief, ...)
- ...

Refining Agents

Advancing agents / relaxing assumptions

- RRL agents with **partial observability**
- RRL agents with **stylised function approximations**
- RRL agents with **finite memory batch** > noisy dynamics
- RRL agents with **different learning paradigms** (imitation, belief, ...)
- ...

→ Taxonomy? 

Refining Evaluations & Analysis

What are we looking for? 

Refining Evaluations & Analysis

What are we looking for? 

- Convergence to equilibria -> multi-stability -> other dynamical regimes (periodic orbits, chaos)

Refining Evaluations & Analysis

What are we looking for? 

- Convergence to equilibria -> multi-stability -> other dynamical regimes (periodic orbits, chaos)
 - **Principles** for their dependence on agents' hyper-parameters and environment characteristics

Refining Evaluations & Analysis

What are we looking for? 

- Convergence to equilibria -> multi-stability -> other dynamical regimes (periodic orbits, chaos)
→ **Principles** for their dependence
on agents' hyper-parameters and environment characteristics
- What else?

Refining Evaluations & Analysis

What are we looking for? 

- Convergence to equilibria -> multi-stability -> other dynamical regimes (periodic orbits, chaos)
→ **Principles** for their dependence
on agents' hyper-parameters and environment characteristics
- What else?

Refining Evaluations & Analysis

What are we looking for? 🔎

- Convergence to equilibria -> multi-stability -> other dynamical regimes (periodic orbits, chaos)
 - ➡ **Principles** for their dependence on agents' hyper-parameters and environment characteristics
- What else?
- ➡ Complex environments <=> empirical (stochastic) game theoretic analysis (Tuyls et al., 2018)

Refining Evaluations & Analysis

What are we looking for? 

- Convergence to equilibria -> multi-stability -> other dynamical regimes (periodic orbits, chaos)
 - **Principles** for their dependence on agents' hyper-parameters and environment characteristics
 - What else?
-
- → Complex environments \Leftrightarrow empirical (stochastic) game theoretic analysis (Tuyls et al., 2018)
 - → Large-agent limit \Leftrightarrow Mean-field games (Elie et al., 2020)

Refining Evaluations & Analysis

What are we looking for? 🔎

- Convergence to equilibria -> multi-stability -> other dynamical regimes (periodic orbits, chaos)
 - **Principles** for their dependence on agents' hyper-parameters and environment characteristics
- What else?
- → Complex environments \Leftrightarrow empirical (stochastic) game theoretic analysis (Tuyls et al., 2018)
- → Large-agent limit \Leftrightarrow Mean-field games (Elie et al., 2020)
- → ... other connections?

Conclusions

Collective Learning Dynamics - a useful semi-formal method
for helping with MARL's challenges

Conclusions

Collective Learning Dynamics - a useful semi-formal method
for helping with MARL's challenges => unified treatment

Conclusions

Collective Learning Dynamics - a useful semi-formal method
for helping with MARL's challenges => unified treatment

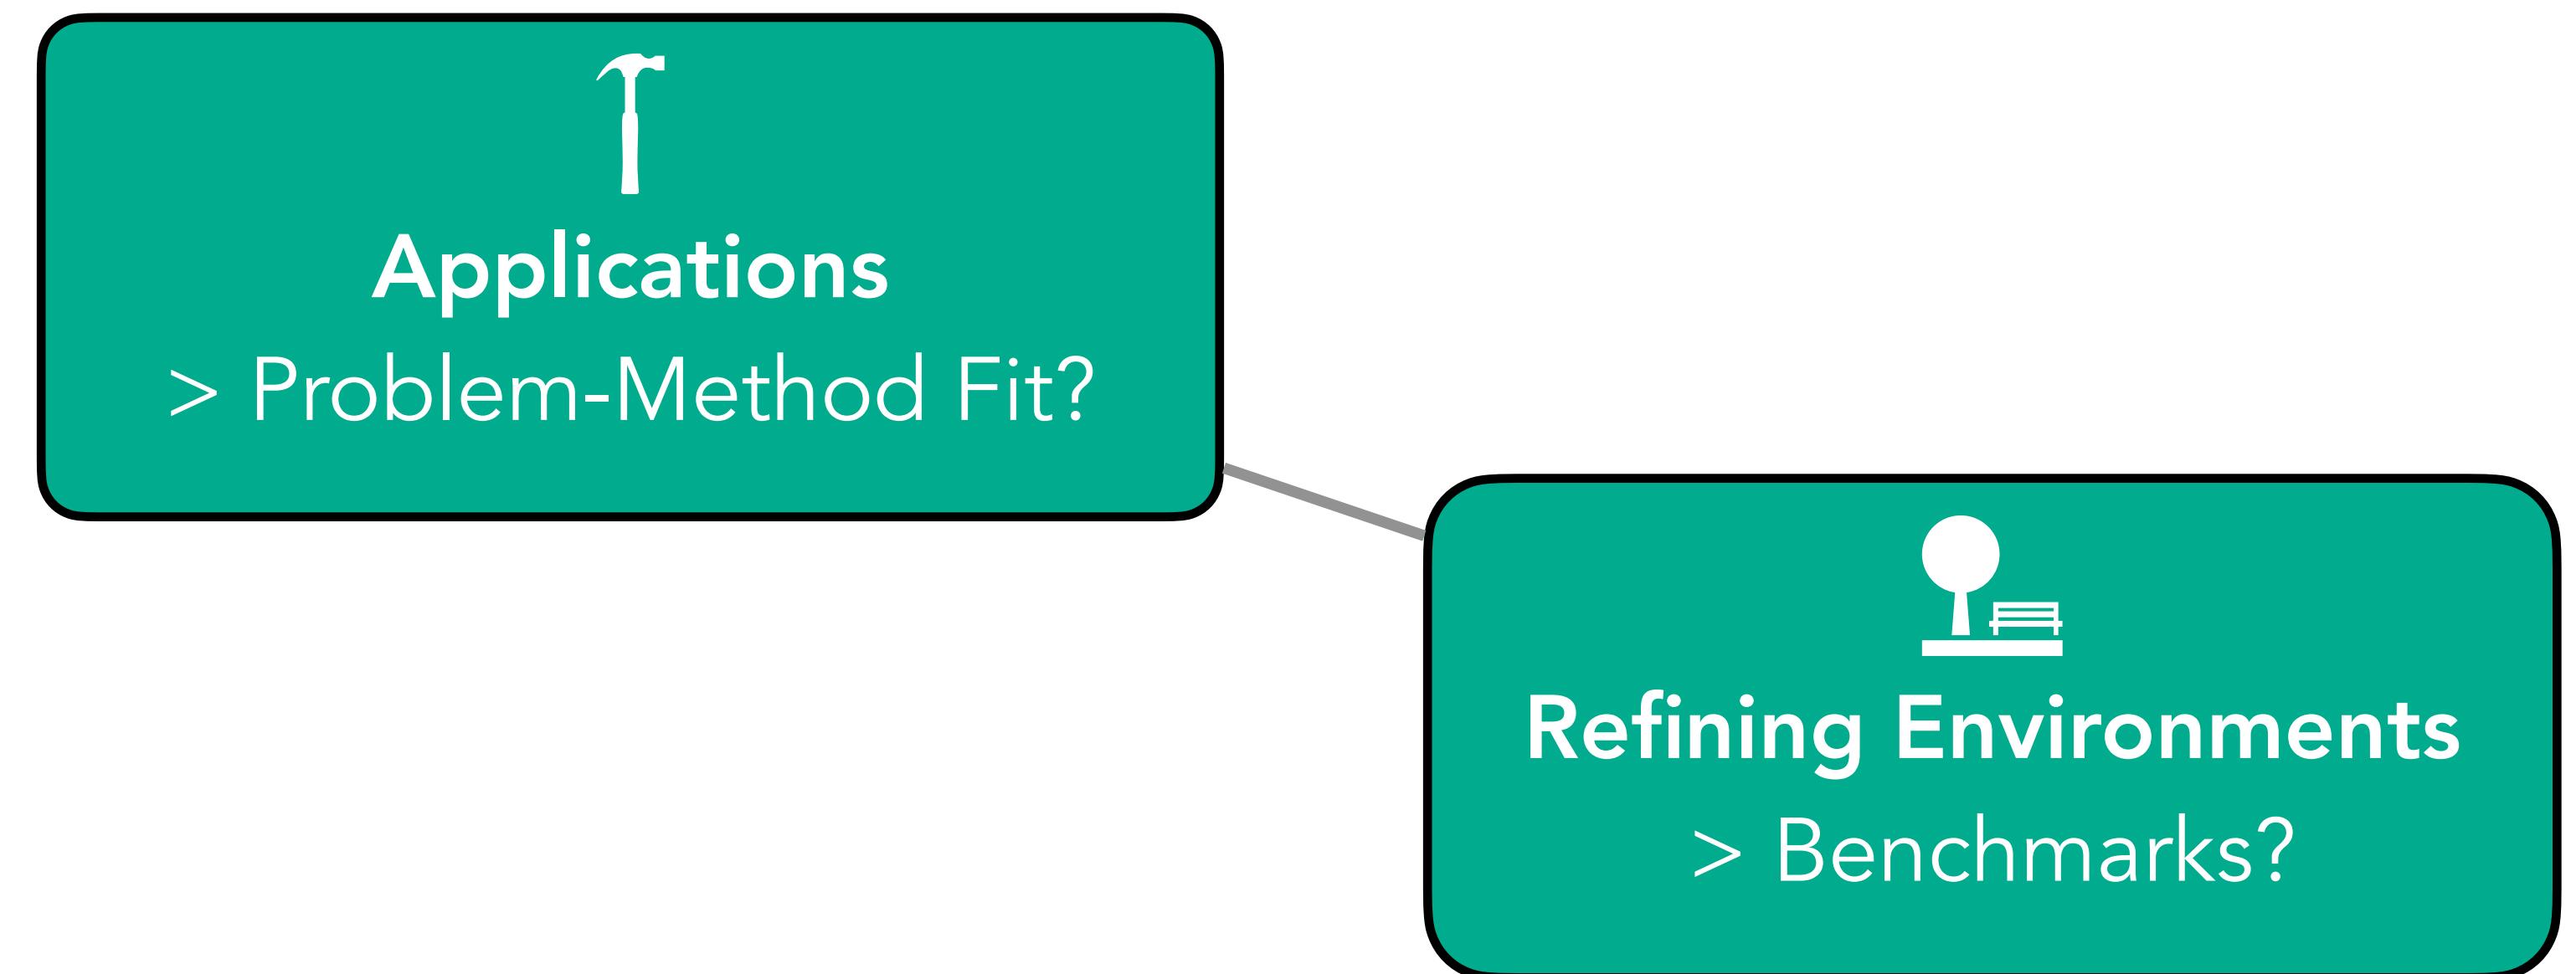


Applications

> Problem-Method Fit?

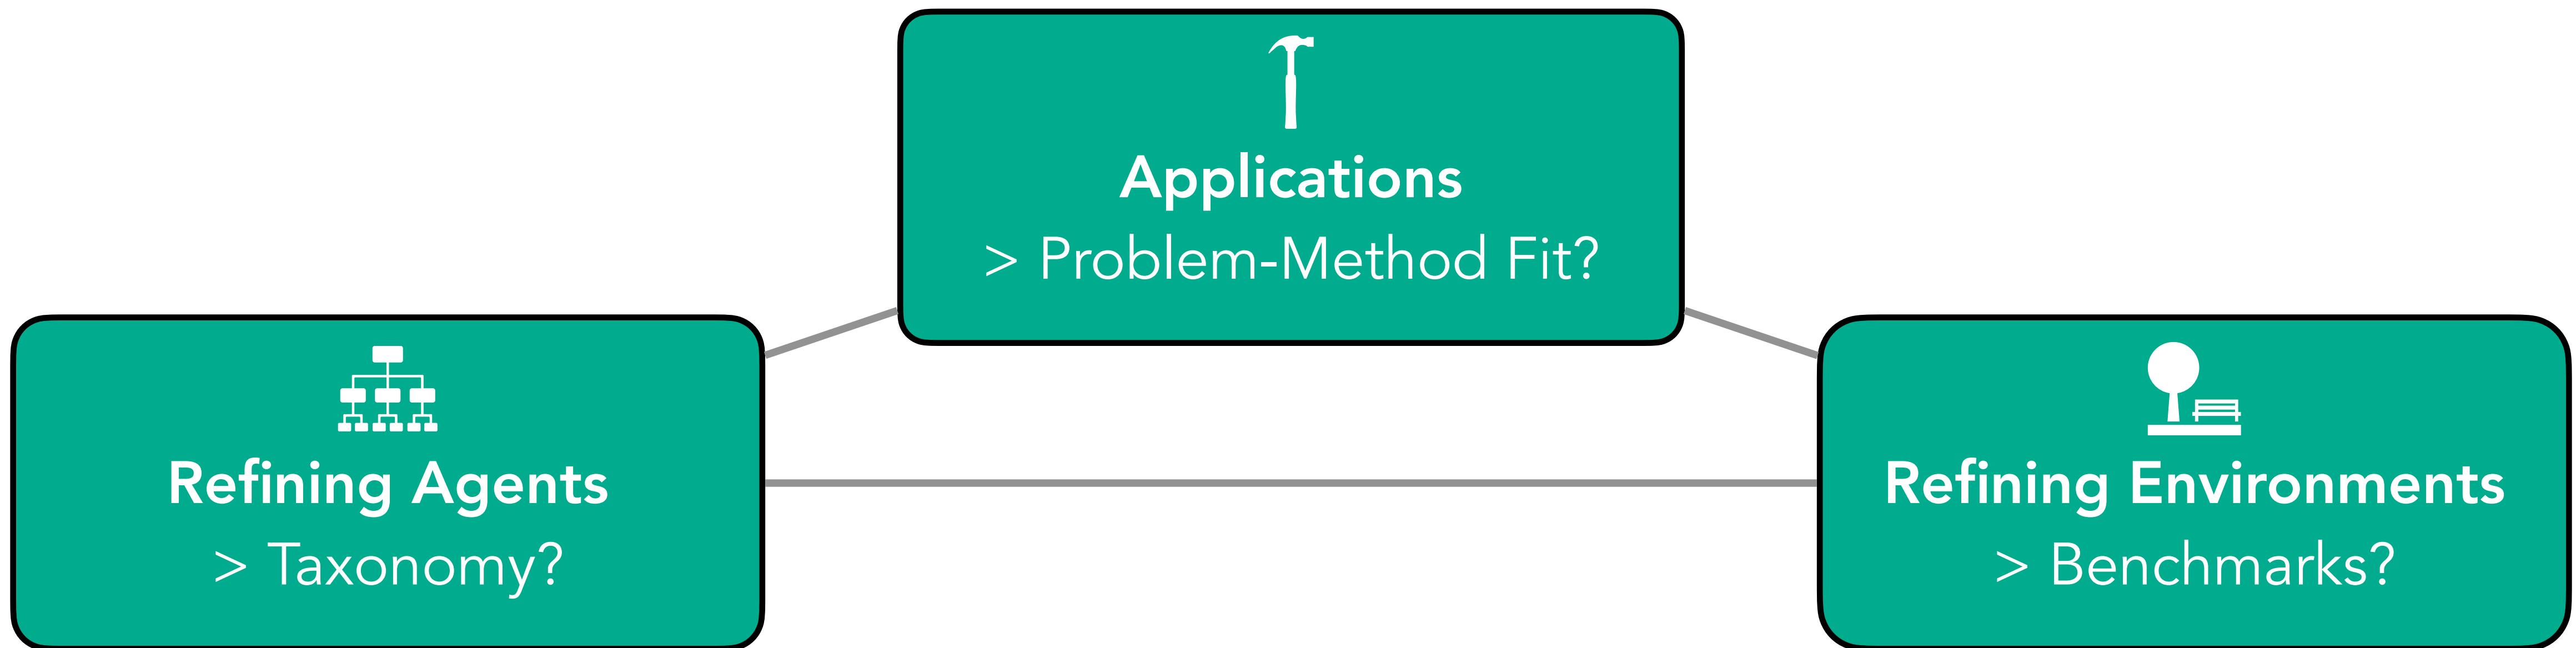
Conclusions

Collective Learning Dynamics - a useful semi-formal method
for helping with MARL's challenges => unified treatment



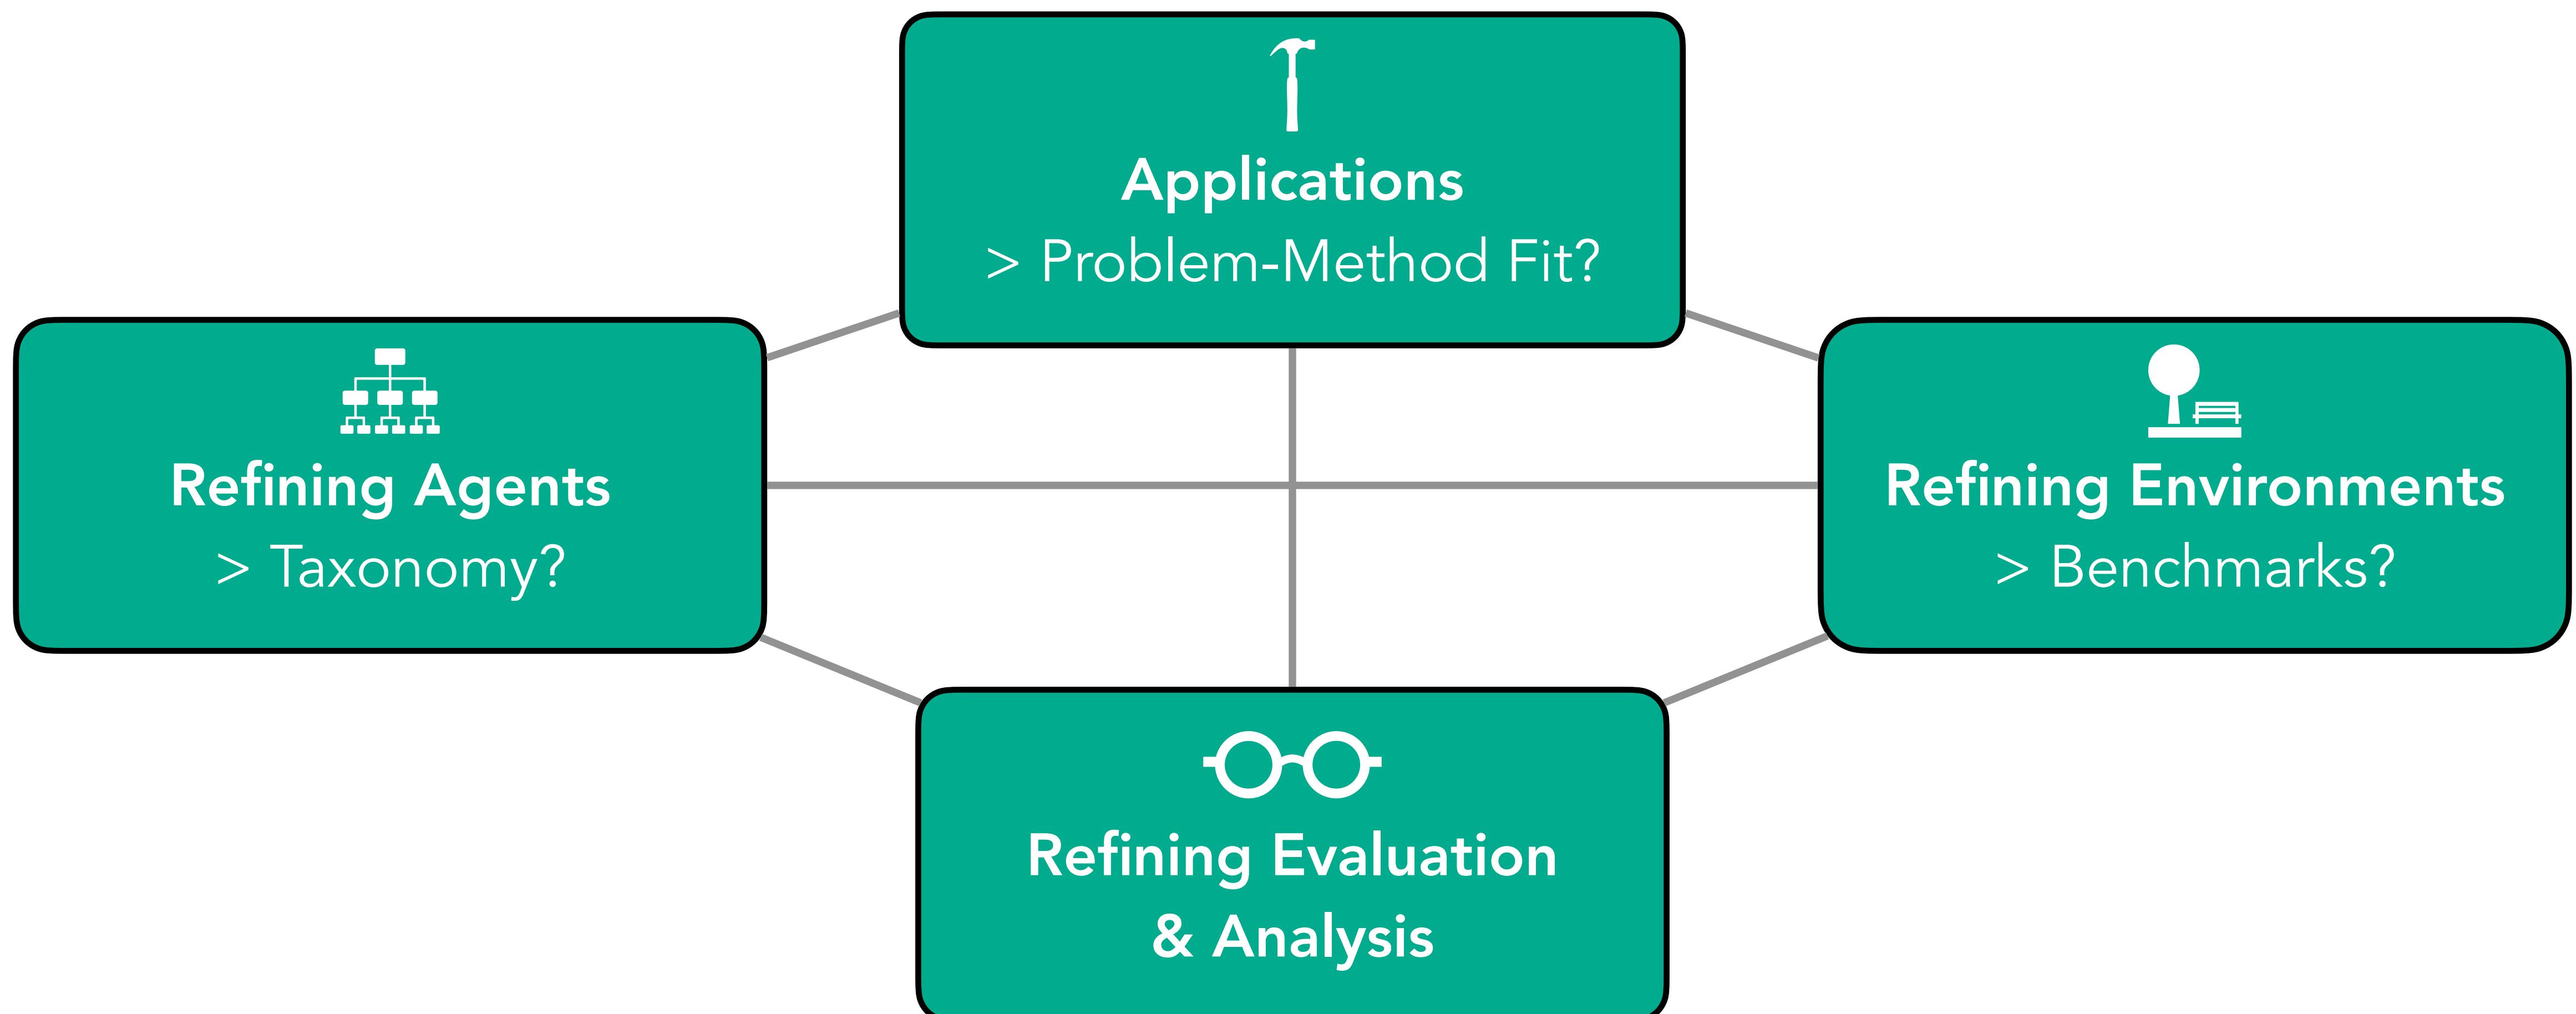
Conclusions

Collective Learning Dynamics - a useful semi-formal method
for helping with MARL's challenges => unified treatment



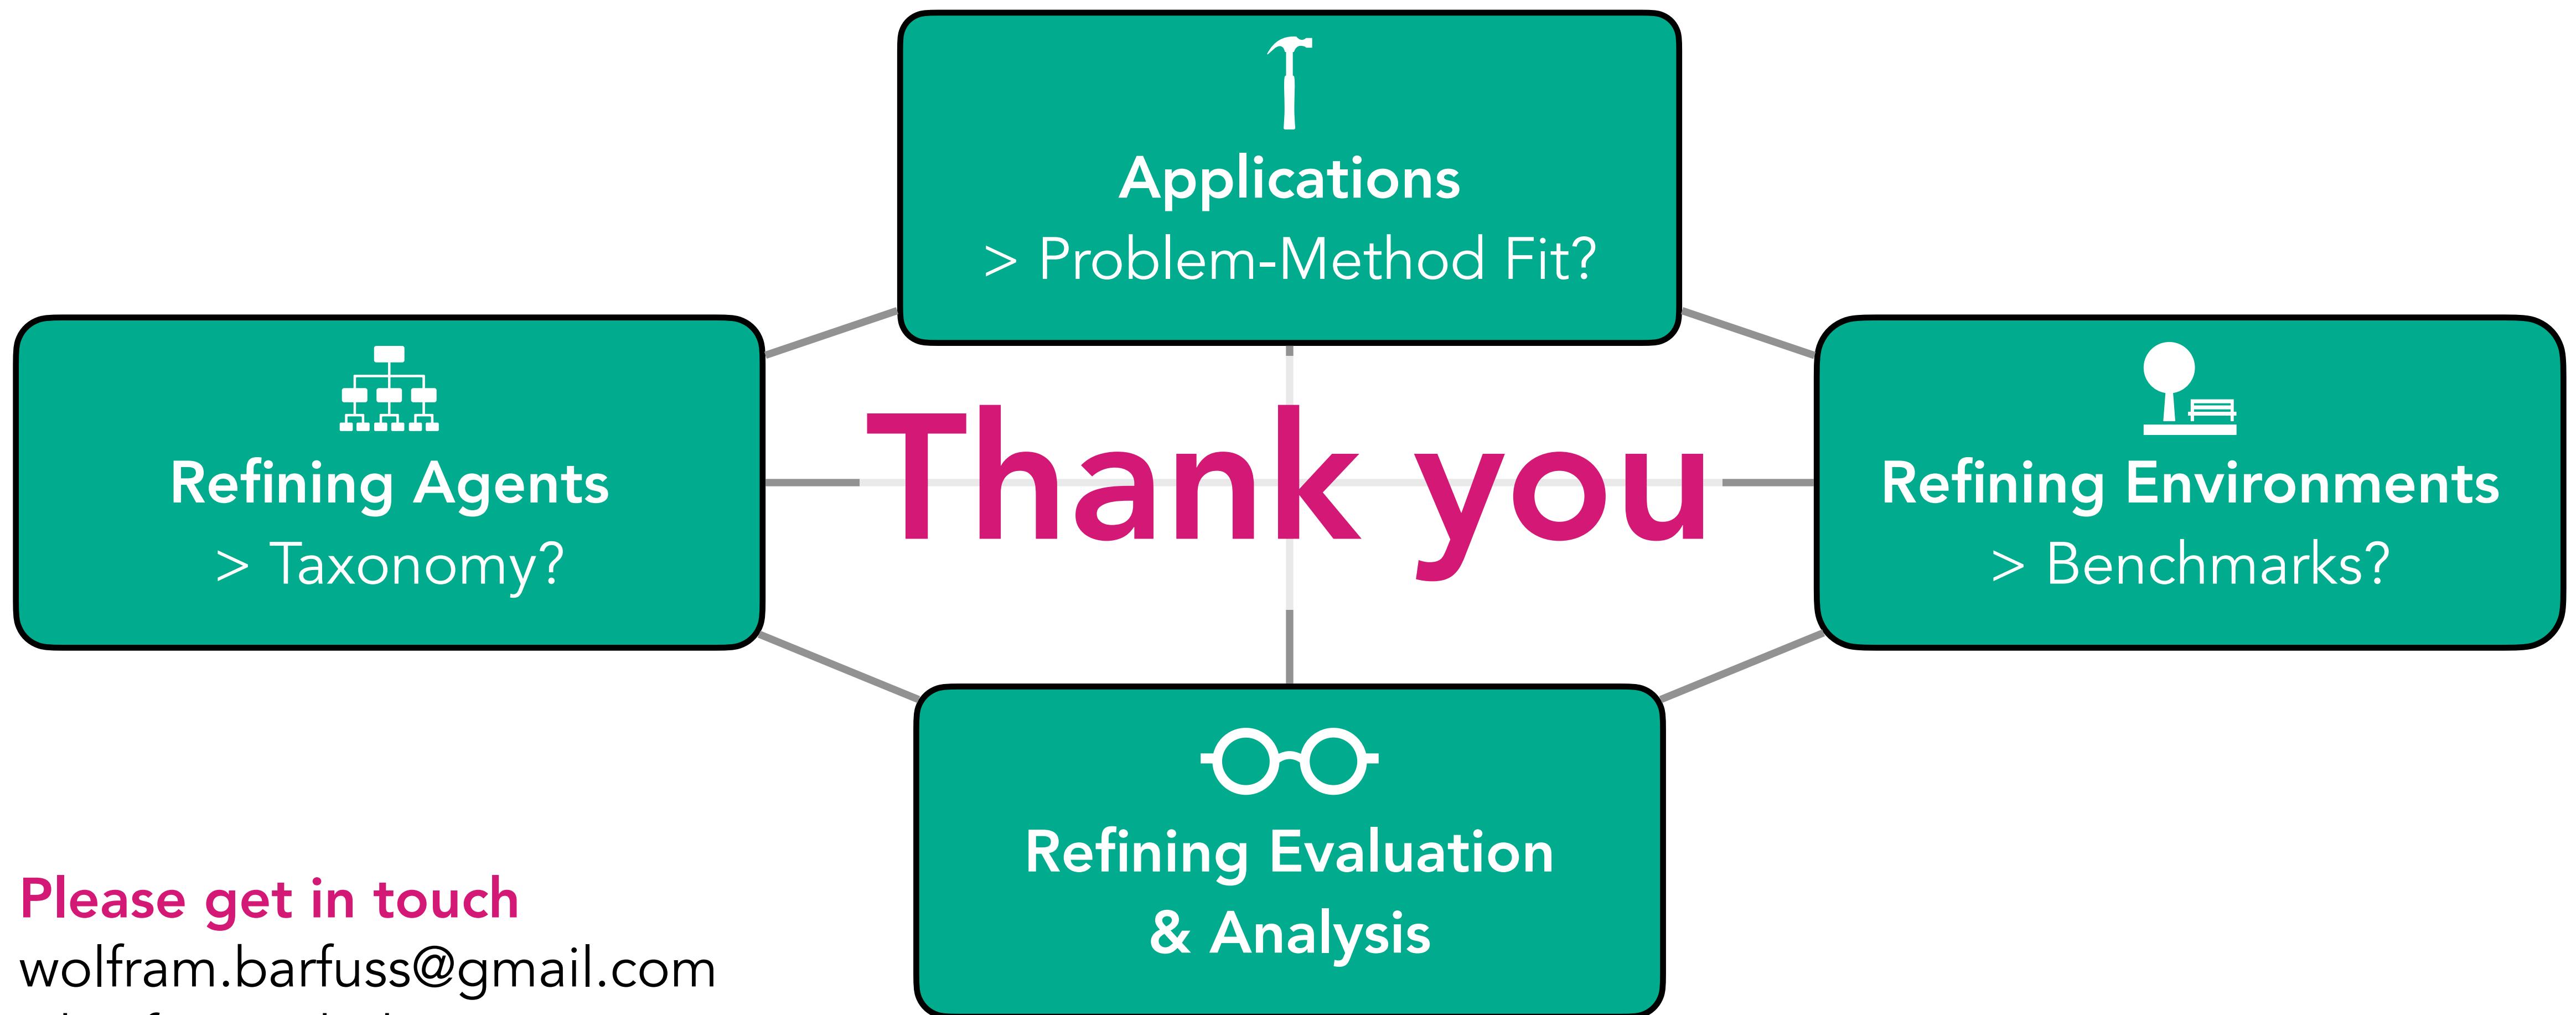
Conclusions

Collective Learning Dynamics - a useful semi-formal method
for helping with MARL's challenges => unified treatment



Conclusions

Collective Learning Dynamics - a useful semi-formal method
for helping with MARL's challenges => unified treatment



Please get in touch

wolfram.barfuss@gmail.com
wbarfuss.github.io

References

- Barfuss, W., Donges, J.F., Vasconcelos, V.V., Kurths, J. and Levin, S.A., 2020. Caring for the future can turn tragedy into comedy for long-term collective action under risk of collapse. *Proceedings of the National Academy of Sciences*, 117(23), pp.12915-12922.
- Bloembergen, D., Tuyls, K., Hennes, D. and Kaisers, M., 2015. Evolutionary dynamics of multi-agent learning: A survey. *Journal of Artificial Intelligence Research*, 53, pp.659-697.
- Elie, R., Pérolat, J., Laurière, M., Geist, M. and Pietquin, O., 2020, April. On the convergence of model free learning in mean field games. In *Proceedings of the AAAI Conference on Artificial Intelligence* (Vol. 34, No. 05, pp. 7143-7150).
- Fudenberg, D., Drew, F., Levine, D.K. and Levine, D.K., 1998. *The theory of learning in games* (Vol. 2). MIT press.
- Han, T.A., Pereira, L.M., Santos, F.C. and Lenaerts, T., 2020. To Regulate or Not: A Social Dynamics Analysis of an Idealised AI Race. *Journal of Artificial Intelligence Research*, 69, pp.881-921.
- Hofbauer, J. and Sigmund, K., 1998. *Evolutionary games and population dynamics*. Cambridge university press.
- Leike, J., Martic, M., Krakovna, V., Ortega, P.A., Everitt, T., Lefrancq, A., Orseau, L. and Legg, S., 2017. AI safety gridworlds. *arXiv preprint arXiv:1711.09883*.
- Osband, I., Doron, Y., Hessel, M., Aslanides, J., Sezener, E., Saraiva, A., McKinney, K., Lattimore, T., Szepesvari, C., Singh, S. and Van Roy, B., 2019. Behaviour suite for reinforcement learning. *arXiv preprint arXiv:1908.03568*.
- Pérolat, J., Leibo, J.Z., Zambaldi, V.F., Beattie, C., Tuyls, K. and Graepel, T., 2017, January. A multi-agent reinforcement learning model of common-pool resource appropriation. In *NIPS*.
- Sato, Y. and Crutchfield, J.P., 2003. Coupled replicator equations for the dynamics of learning in multiagent systems. *Physical Review E*, 67(1), p.015206.
- Tuyls, K., Perolat, J., Lanctot, M., Leibo, J.Z. and Graepel, T., 2018. A generalised method for empirical game theoretic analysis. *arXiv preprint arXiv:1803.06376*.